

ECE 568: Computer Security

Part 1: Course Overview and Review

Courtney Gibson, P.Eng.



UNIVERSITY OF
TORONTO

Part 1

Course Overview and Review

- Introduction
- Why should you take this course?
- Course Overview
- Security Fundamentals
- Review (Comp Org and Operating Systems)



Hello
my name is

COURTNEY

Introduction

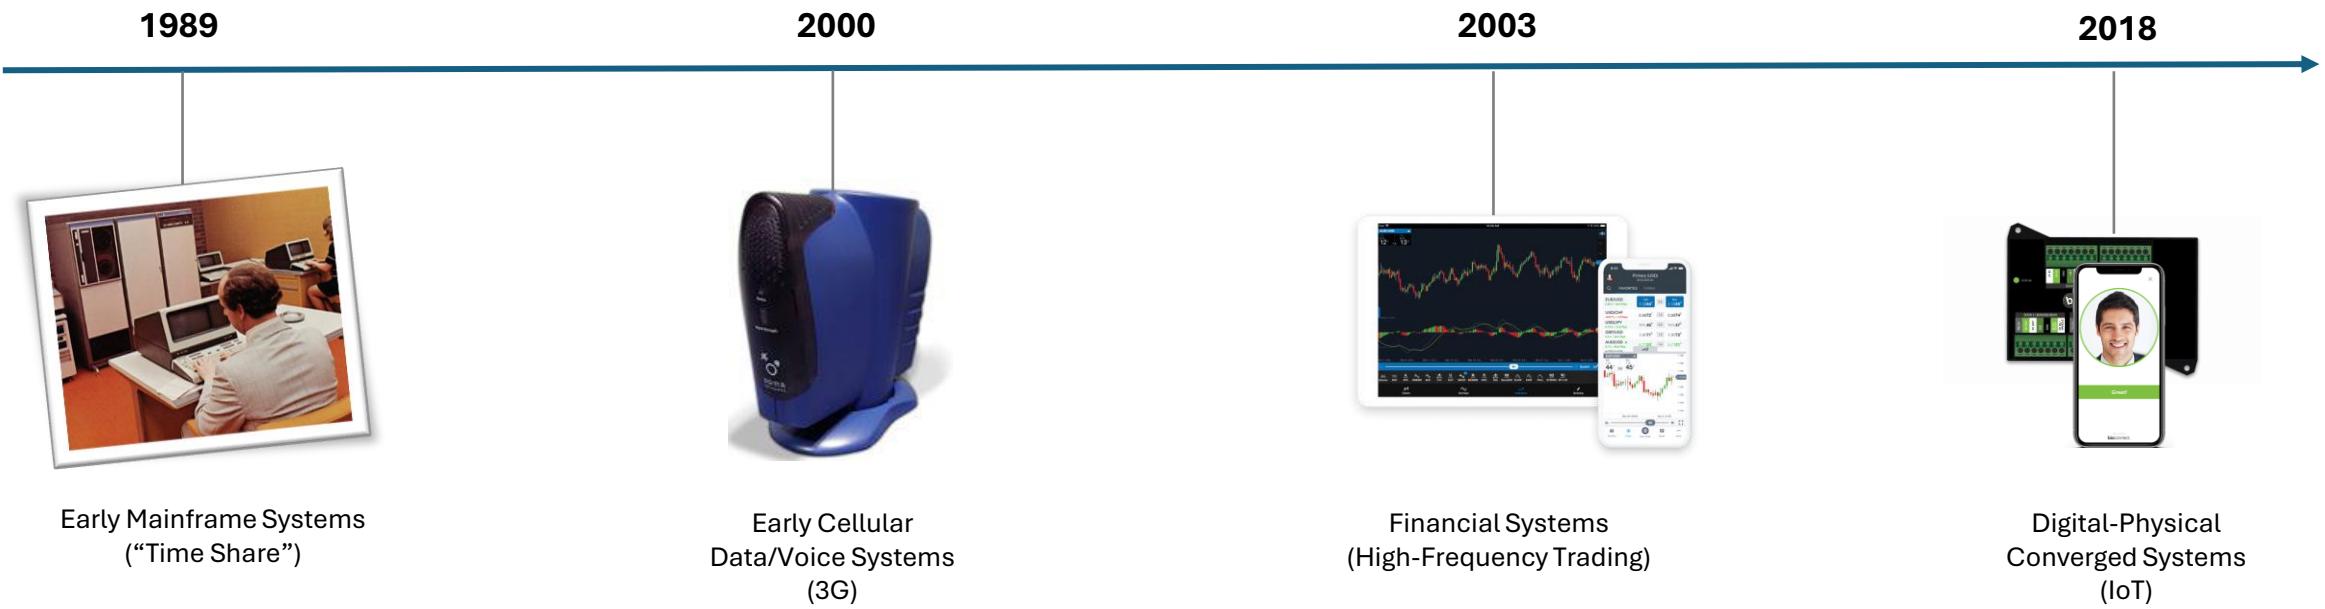


UNIVERSITY OF
TORONTO



Courtney Gibson

courtney.gibson@utoronto.ca



Early Mainframe Systems
("Time Share")

Early Cellular
Data/Voice Systems
(3G)

Financial Systems
(High-Frequency Trading)

Digital-Physical
Converged Systems
(IoT)



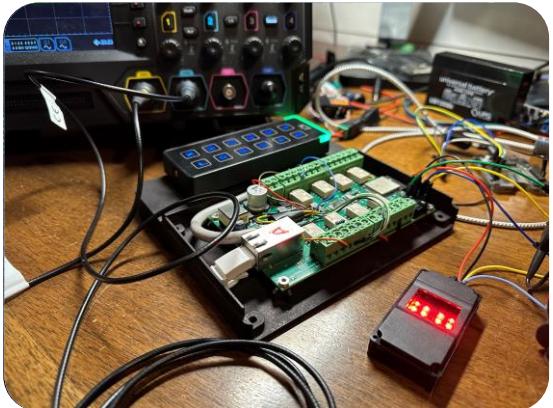
Courtney Gibson

courtney.gibson@utoronto.ca

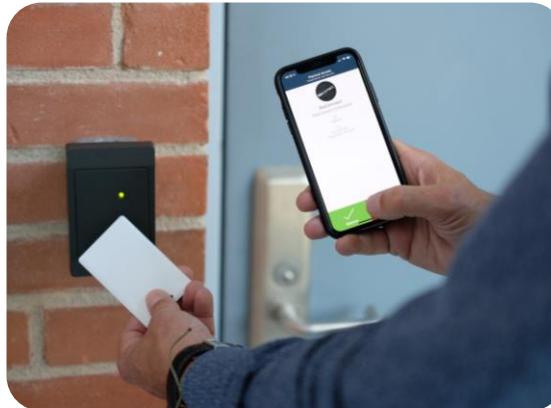
bioconnect.

Chief Technology Officer
Chief Information Security Officer

IoT Design



Biometrics, RFID and Access Control



Compliance





Courtney Gibson

courtney.gibson@utoronto.ca



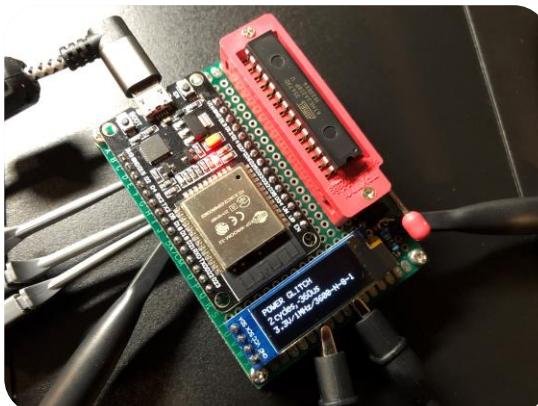
Adjunct Lecturer, Electrical and Computer Engineering
Instructor, Rotman School of Management

Cybersecurity

Buffer Overflow

After the return address is changed, when the function returns, it will return to the changed return address.

An attacker can use this vulnerability to **hijack** the program (i.e., alter the program instruction stream).



Cyber-Risk Management

Cyber-Physical Convergence

Organizations should adopt specific policies to manage their risks from integrated cyber-physical systems.



- Security:** Product evaluations should include a robust validation of end-to-end security, including:
 - How data is protected from theft or tampering
 - How devices and servers are protected from unauthorized changes to software or settings
 - Documented assurance there are no “back doors” or hidden “service accounts”

- Patching:** Critical vulnerabilities should be quickly and automatically patched
 - Centralized alerting for unpatched devices

- Zero-Trust:** Access to other resources within the organization should require authentication
 - Restrict access to only the minimum resources required

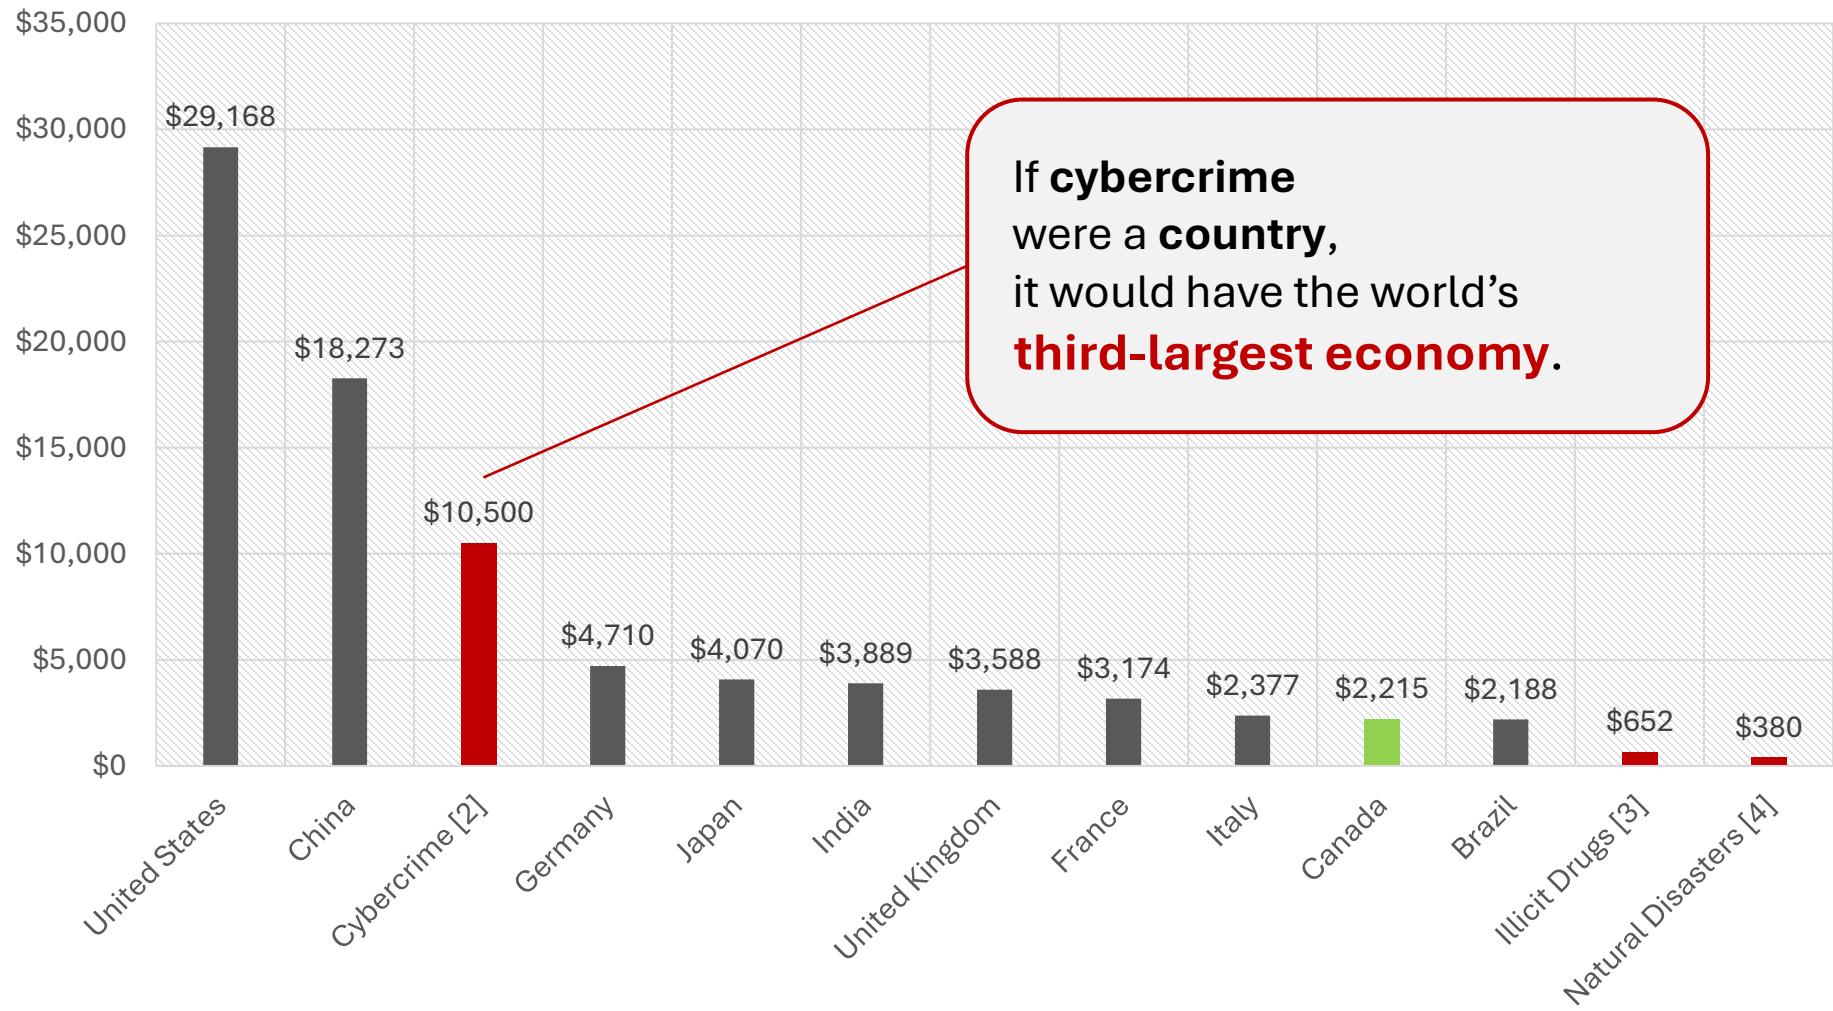
Please Ask Questions!





Why Should You Take This Course?

Top-10 Countries by GDP (2024¹, \$B USD)



¹ IMF, *World Economic Outlook Database (October 2024)*

² Cybersecurity Ventures, *Boardroom Cybersecurity Report 2024*

³ GFI, *Transnational Crime and the Developing World (2017)*

⁴ AON, *Climate and Catastrophe Insight 2024*

DEMO

You have spent a **lot of time** over the past few years learning the skills you need to analyze problems and solve them efficiently.

“Write a program to sort of a list of up to 100 numbers, as efficiently as possible.”

You have spent **very little time** (if any) thinking about how an attacker might **break your work...** and **what might happen** if they succeed.

Lab Assignment #7: Inheritance

1. Objectives

The objective of this assignment is to provide you with practice on the use of inheritance in C++ programming. This will be done in the context of re-implementing the simple student-marks database of Assignment 5 to allow the storage and retrieval of records of any type, not only of type `studentRecord`.

2. Problem Statement

In this assignment, you will implement a simple array-based database to store and retrieve records. In the first part of the assignment, you will implement two classes: `Record` and `DB`. The `Record` class will serve as a base class from which other types of record classes can be derived. The `DB` class will be used to create a database of `Record` objects. In the second part of the assignment, you will design and implement the class `studentRecord`, which is derived from the class `Record`. You will test your implementation with the `Driver` you wrote for Assignment 5. However, your implementation of `Record` and `DB` must work for any class that is derived from `Record`, even without any knowledge on your part of what the derived class does.

2.1 The Record Class

~~The `Record` class has fields to represent the number (key) of an individual. It also has the~~

~~class Record {~~

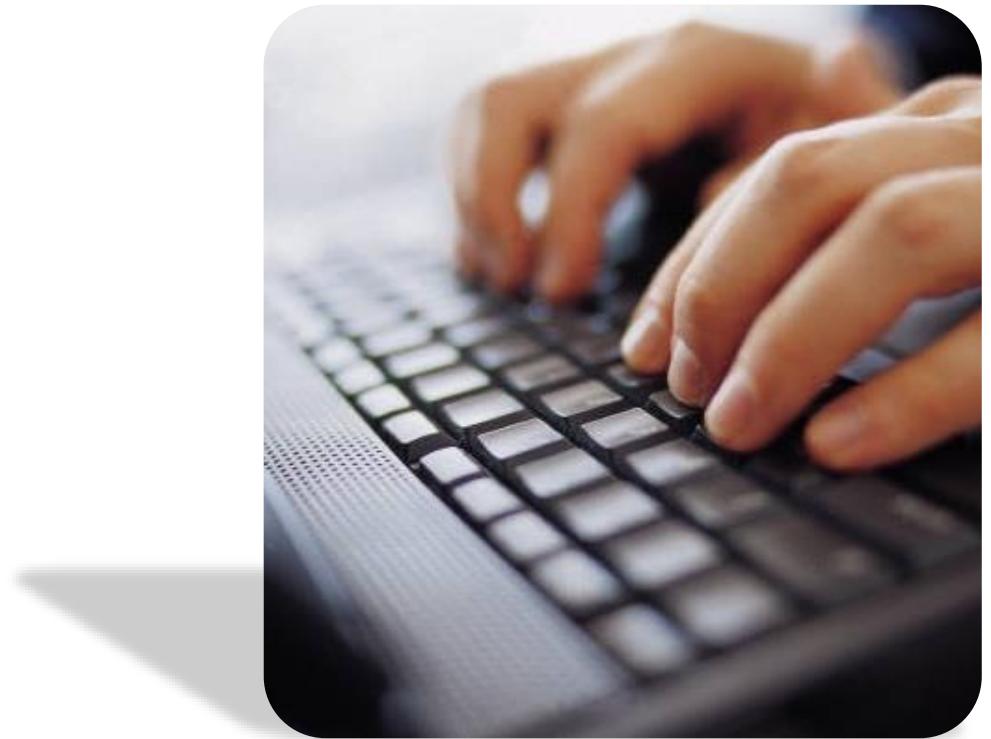
~~public:~~

~~Record();~~

Motivation: Users of Technology

Software systems (and their associated cybersecurity risks) are ubiquitous in our daily lives.

- Protect your system
- Protect your data
- Protect your identity

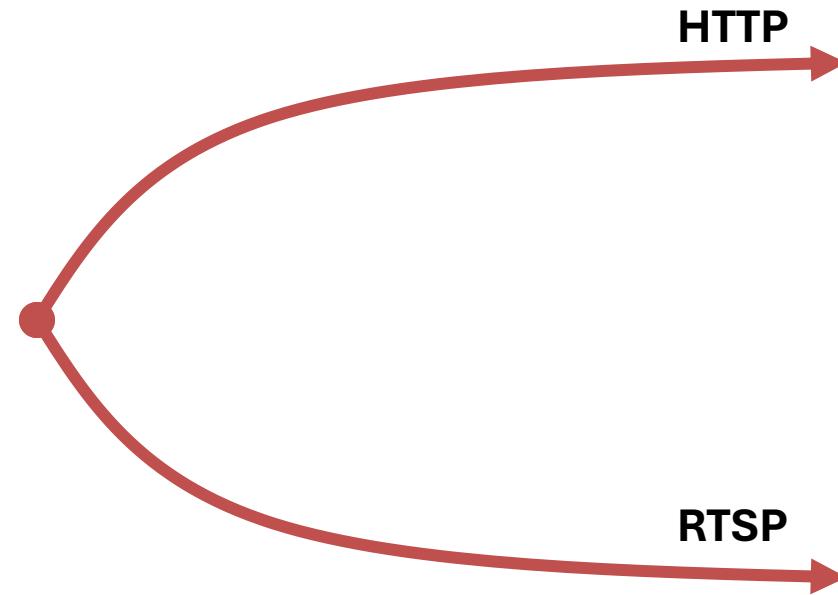


Somebody's Watching: Hackers Breach Ring Home Security Cameras

Unnerved owners of the devices reported recent hacks in four states. The company reminded customers not to recycle passwords and user names.



A family in Mississippi said a man hacked into a Ring home security camera in a bedroom shared by their daughters. Ashley LeMay

A screenshot of a login interface. It features two input fields: "User Name" with a user icon and "Password" with a lock icon. Below the fields is a large red "Login" button.

RESEARCH NEWS ABOUT Q

munk school OF GLOBAL AFFAIRS & PUBLIC POLICY UNIVERSITY OF TORONTO

THECITIZENLAB

Research > Transparency and Accountability

Finding You The Network Effect of Telecommunications Vulnerabilities for Location Disclosure

By Gary Miller and Christopher Parsons October 26, 2024

[Download this report](#)

Table of Contents

- 1. Introductions
- 2. Roaming, SIMs, and Services 101
- 3. Geolocation Attacks Against Telecommunications Networks
- 4. Case Studies and Statistics
- 5. Incentives Enabling Geolocation Attacks
- 6. Geolocation Tracking in 5G Networks and Unimplemented Features
- 6. Conclusion

Introduction

The information collected by, and stored within, these networks and reveal our behaviours, dining habits, sleeping patterns, and where we live.

Ron Deibert and Noura Aljizawi join CBC chief correspondent Adrienne Arsenault to discuss the growing issue of foreign governments' threats and intimidation targeting exiled women activists and dissidents living in Canada.

In this segment, women activists from Iran and China recount their experiences of digital transnational repression while living in Canada. They describe how perpetrators use surveillance and intimidation tactics to force them into silence and isolation.

RESEARCH NEWS ABOUT Q

munk school OF GLOBAL AFFAIRS & PUBLIC POLICY UNIVERSITY OF TORONTO

THECITIZENLAB

Research > App Privacy and Controls

The not-so-silent type Vulnerabilities across keyboard apps reveal keystrokes

By Jeffrey Knockel and Zoë Reichert April 23, 2024

讀本報告的繁體中文版摘要 | A mirrored copy of this report is available on

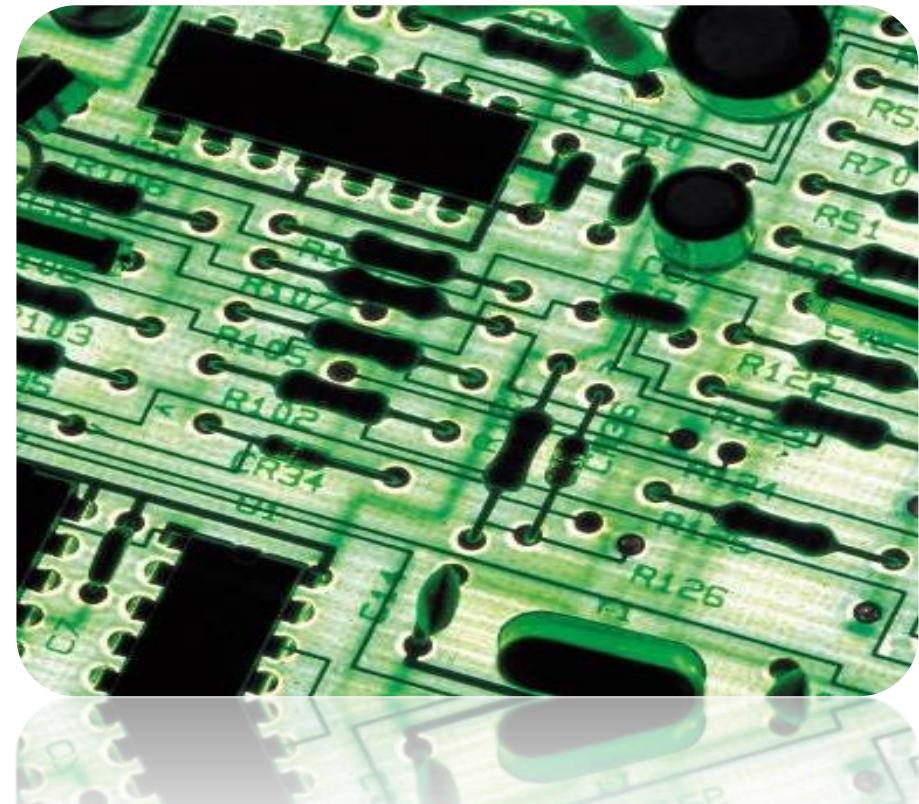
updates to [their keyboard apps](#) and that they keep their mobile operating system at-risk users consider switching from a cloud-based keyboard. [Read the FAQ accompanying this report.](#)

-based pinyin keyboard apps from nine vendors — PO, Samsung, Tencent, Vivo, and Xiaomi — and users' keystrokes for vulnerabilities. abilities in keyboard apps from eight out of the nine most vulnerable to completely reveal the contents the vulnerable apps can be exploited by an

Motivation: Engineer

Security should be a daily consideration in your work.

- Challenge your design assumptions
- Understand attackers and identify risks
- Defensive architecture and implementation



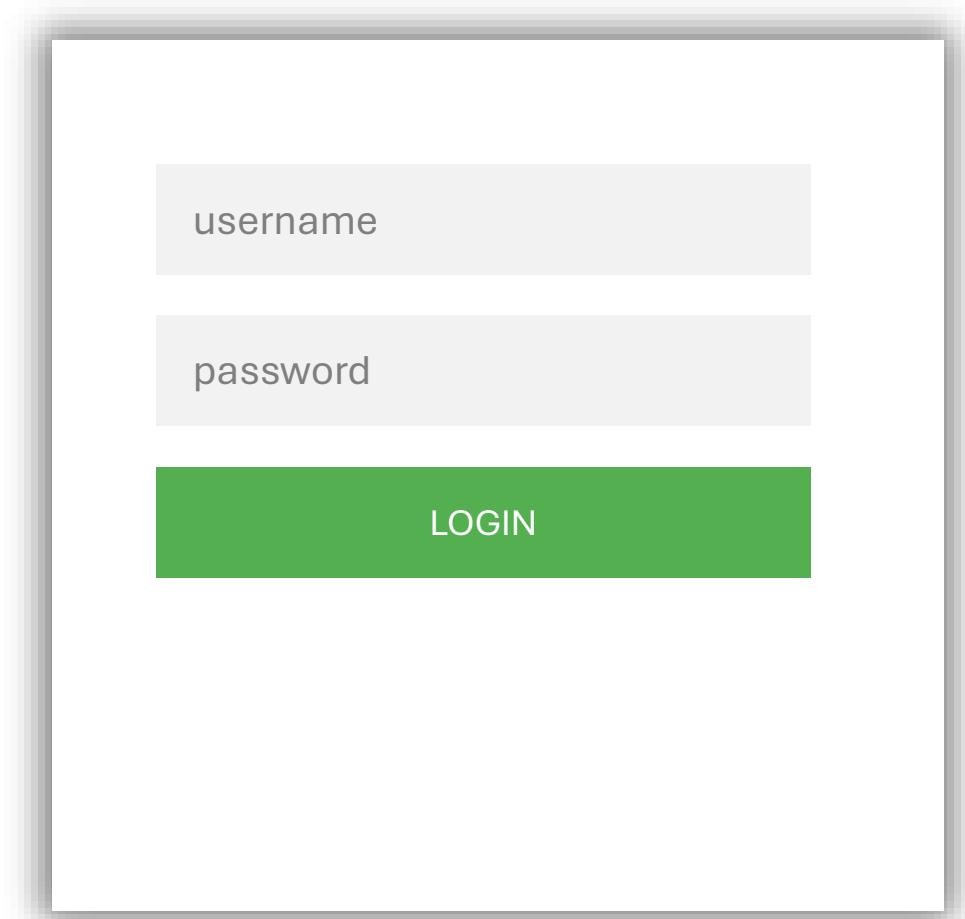
Case Study: Overly-Helpful Login

Background

- New hire at a financial services company was given the task of developing a web form that would allow customers to download their annual tax forms.

Analysis

- System should not have been storing passwords in plaintext!
- Error message leaked too much information to potential attackers.



Case Study: Las Vegas Casino

Background

- In early 2018 a large Las Vegas casino installed a new saltwater fish tank in their lobby; this involved installing an IoT device to permit remote monitoring of the water conditions in the tank

Issue

- The IoT tank monitor had a vulnerability: attackers remotely accessed the device, and then used it as an entry-point into the casino's internal network

Result

- The contents of the casino's "high-roller" database (with PII on thousands of VIP clients) was sent to a cloud server hosted in Finland (~10GB of data)
- Several million dollars in ransom was paid to prevent the release of the data

Analysis

- IoT devices are often overlooked, and avoid the same level of review that other new computers receive





Industrial Control Systems

The Basics

Industrial control systems (ICS) are computers that control the world around you. They're responsible for managing the air conditioning in your office, the turbines at a power plant, the lighting at the theatre or the robots at a factory.

Common Terms

SCADA	Supervisory Control and Data Acquisition
PLC	Programmable Logic Controller
DCS	Distributed Control System

Search Filter

Shodan continuously crawls the Internet and discovers Internet-accessible ICS devices. If you have an enterprise subscription to Shodan you can use the **tag** search filter with a value of **ics** to get a list all ICS on the Internet right now.

[EXPLORE ICS](#)

Modbus

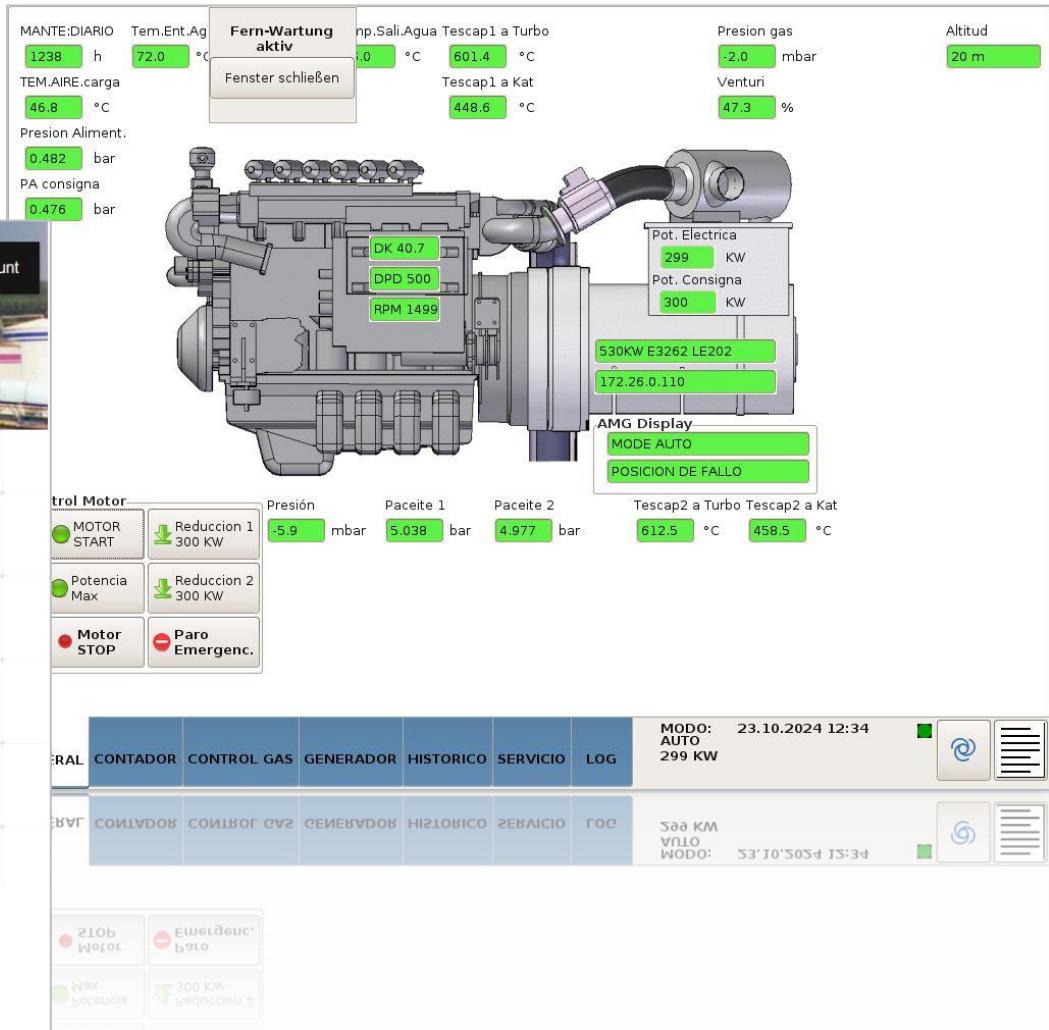
Modbus is a popular protocol for industrial control systems (ICS). It provides easy, raw access to the control system without requiring any authentication.

SIEMENS

S7 (S7 Communication) is a Siemens proprietary protocol that runs between programmable logic controllers (PLCs) of the Siemens S7 family.

dnp

DNP3 (Distributed Network Protocol) is a set of communications protocols used between components in process automation systems. Its main use is in utilities such as electric and water



Motivation: Employee

Laws increasingly hold companies accountable for poor security practice.

- Protect your company
- Protection of confidential data
- Regulatory requirements



Case Study: Equifax

Background

- Multinational consumer credit reporting agency: processes data on over 800M individuals and 88M businesses worldwide. In 2017 they experienced a substantial cyber breach.

Issue

- 2015 internal audit showed significant deficiencies with the patching of critical systems:
 - IT staff did not have a comprehensive asset inventory
 - The criticality of an IT asset was not considered when prioritizing patches
 - The patching process worked on an 'Honour system'
 - The report identified remediation actions; two years later many of the steps had not been completed
- A critical web server security patch was released in March 2017
 - Equifax was aware of the issue; had not updated their Credit Dispute website
 - Hacked over two months after the public was notified and urged to install the patch
- Systems designed to detect data exfiltration were offline: a certificate had expired 9 months prior

Result

- Breach lasted for 76 days
- Private records of over 163M American, British and Canadian citizens were compromised
- Multiple class-action lawsuits and regulatory actions

Vulnerabilities

All ports ▾ Latest ▾

Note: the device may not be impacted by all of these issues. The vulnerabilities are implied based on the software and version.

2024

CVE-2024-40898 falseSSRF in Apache HTTP Server on Windows context, allows to potentially leak HTML malicious requests. Users are recommended to upgrade to version 2.4.60, which fixes this issue.

CVE-2024-38477 falsenull pointer dereference in mod_perl earlier allows an attacker to crash the server. It is recommended to upgrade to version 2.4.60, which fixes this issue.

CVE-2024-38476 falseVulnerability in core of Apache HTTP Server to information disclosure, SSRF or local privilege escalation whose response headers are malicious. It is recommended to upgrade to version 2.4.60, which fixes this issue.

CVE-2024-38474 falseSubstitution encoding issue in mod_perl and earlier allows attacker to execute shell command configuration but not directly reachable URLs meant to only to be executed as CGI. It is recommended to upgrade to version 2.4.60, which fixes this issue. Some substitutions may now fail unless explicitly disabled.

CVE-2024-40898 falseSSRF in Apache HTTP Server on Windows context, allows to potentially leak HTML malicious requests. Users are recommended to upgrade to version 2.4.60, which fixes this issue.

CVE-2024-38477 falsenull pointer dereference in mod_perl earlier allows an attacker to crash the server. It is recommended to upgrade to version 2.4.60, which fixes this issue.

CVE-2024-38476 falseVulnerability in core of Apache HTTP Server to information disclosure, SSRF or local privilege escalation whose response headers are malicious. It is recommended to upgrade to version 2.4.60, which fixes this issue.

CVE-2024-38474 falseSubstitution encoding issue in mod_perl and earlier allows attacker to execute shell command configuration but not directly reachable URLs meant to only to be executed as CGI. It is recommended to upgrade to version 2.4.60, which fixes this issue. Some substitutions may now fail unless explicitly disabled.

https://www.shodan.io

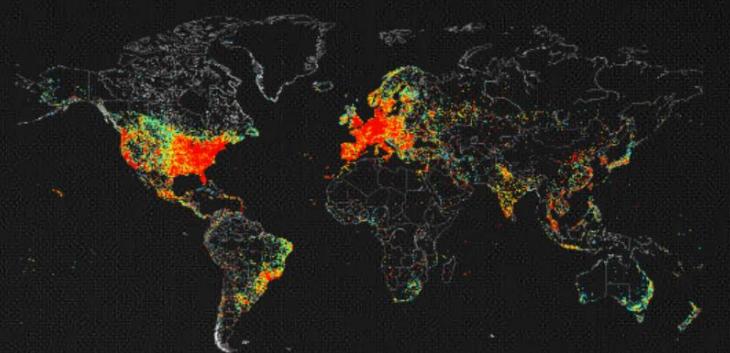
Shodan Maps Images Monitor Developer More...

SHODAN Explore Pricing Search

Search Engine for the Internet of Everything

Shodan is the world's first search engine for Internet-connected devices. Discover how Internet intelligence can help you make better decisions.

SIGN UP NOW



Case Study: Colonial Pipeline

Background

- Largest fuel operator in the United States
- Transports refined gasoline and jet fuel from Houston, across the southeastern United States
 - Approximately 45% of all fuel consumed on the US east coast arrives via their pipeline system

Issue

- May 2021: Ransomware attack took down their billing system. Unable to bill customers, Colonial decided to shut down their pipeline.
 - Fuel shortages led to panic buying: over 11,000 gas stations ran out of fuel in US east coast; impact to airlines at several airports
 - Federal government declared a state of emergency

Result

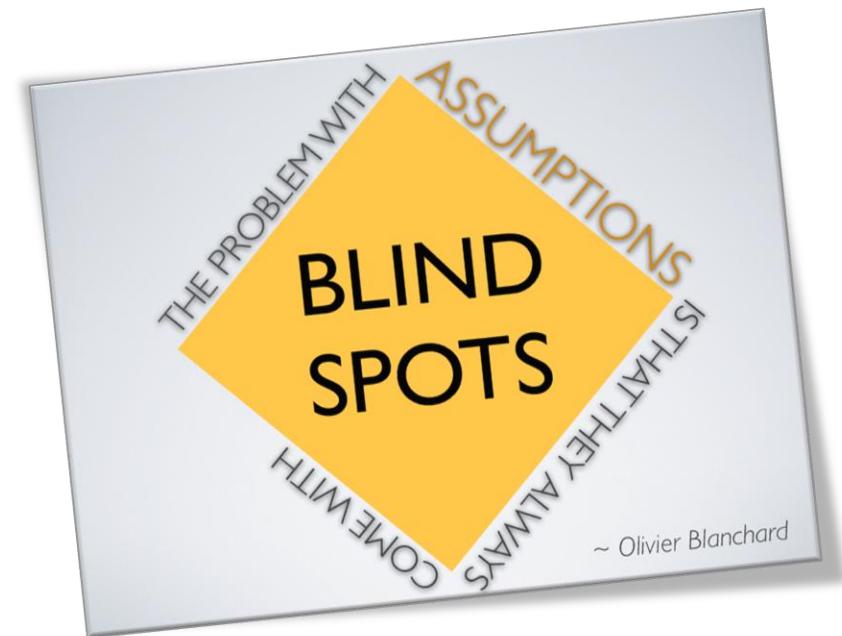
- Paid a ransom of \$4.4M in Bitcoin
 - Attackers stole ~100GB of data day before malware attack; threatened to release it if the ransom wasn't paid
 - FBI recovered 84% of payment – but drop in Bitcoin meant the recovered funds had lost half their value
- A data-recovery tool was provided by the attackers – however, the tool was so slow that Colonial's own BCP procedures were more-effective in restoring their systems. Led to a six-day shutdown and significant operating losses.

My two goals for you in this course are to help you:

1. Get better at **questioning your assumptions**.
 2. Better-understand **what tools you can trust, and when**.
-

We will be spending this term looking closely at the limits of trust you can put in your tools:

- Languages, Compilers
- Hardware
- Networking
- Storage





Microsoft Partner
Cloud Productivity
Silver Partner
Midmarket Solution Provider

KELSER Ten Immutable Laws of Security (Version 2.0)

Here at the Microsoft Security Response Center, we investigate thousands of security reports every year. In some cases, we find that a report describes a bona fide security vulnerability resulting from an issue in one of our products; when this happens, we develop a corrective update as quickly as possible. In other cases, the reported problems simply result from a mistake someone made in using the product, or our investigation finds a problem with the product that, while troublesome for users, does not expose them to a security vulnerability. But many fall in between. They are genuine security problems, but the problems don't result from product flaws. Over the years, we've developed a list of issues like these that we call the 10 Immutable Laws of Security.

Don't hold your breath waiting for an update that will protect you from the issues we'll discuss below. It isn't possible for Microsoft—or any software vendor—to "fix" them, because they result from the way computers work. But don't abandon all hope yet. Sound judgment is the key to protecting yourself against these pitfalls, and if you keep them in mind, you can significantly improve the security of your computers, whether they sit on your desk, travel in your pocket, or exist in a virtual cloud. (Throughout this list we'll use "computer" to mean all of those objects, by the way.)

The 10 Immutable Laws

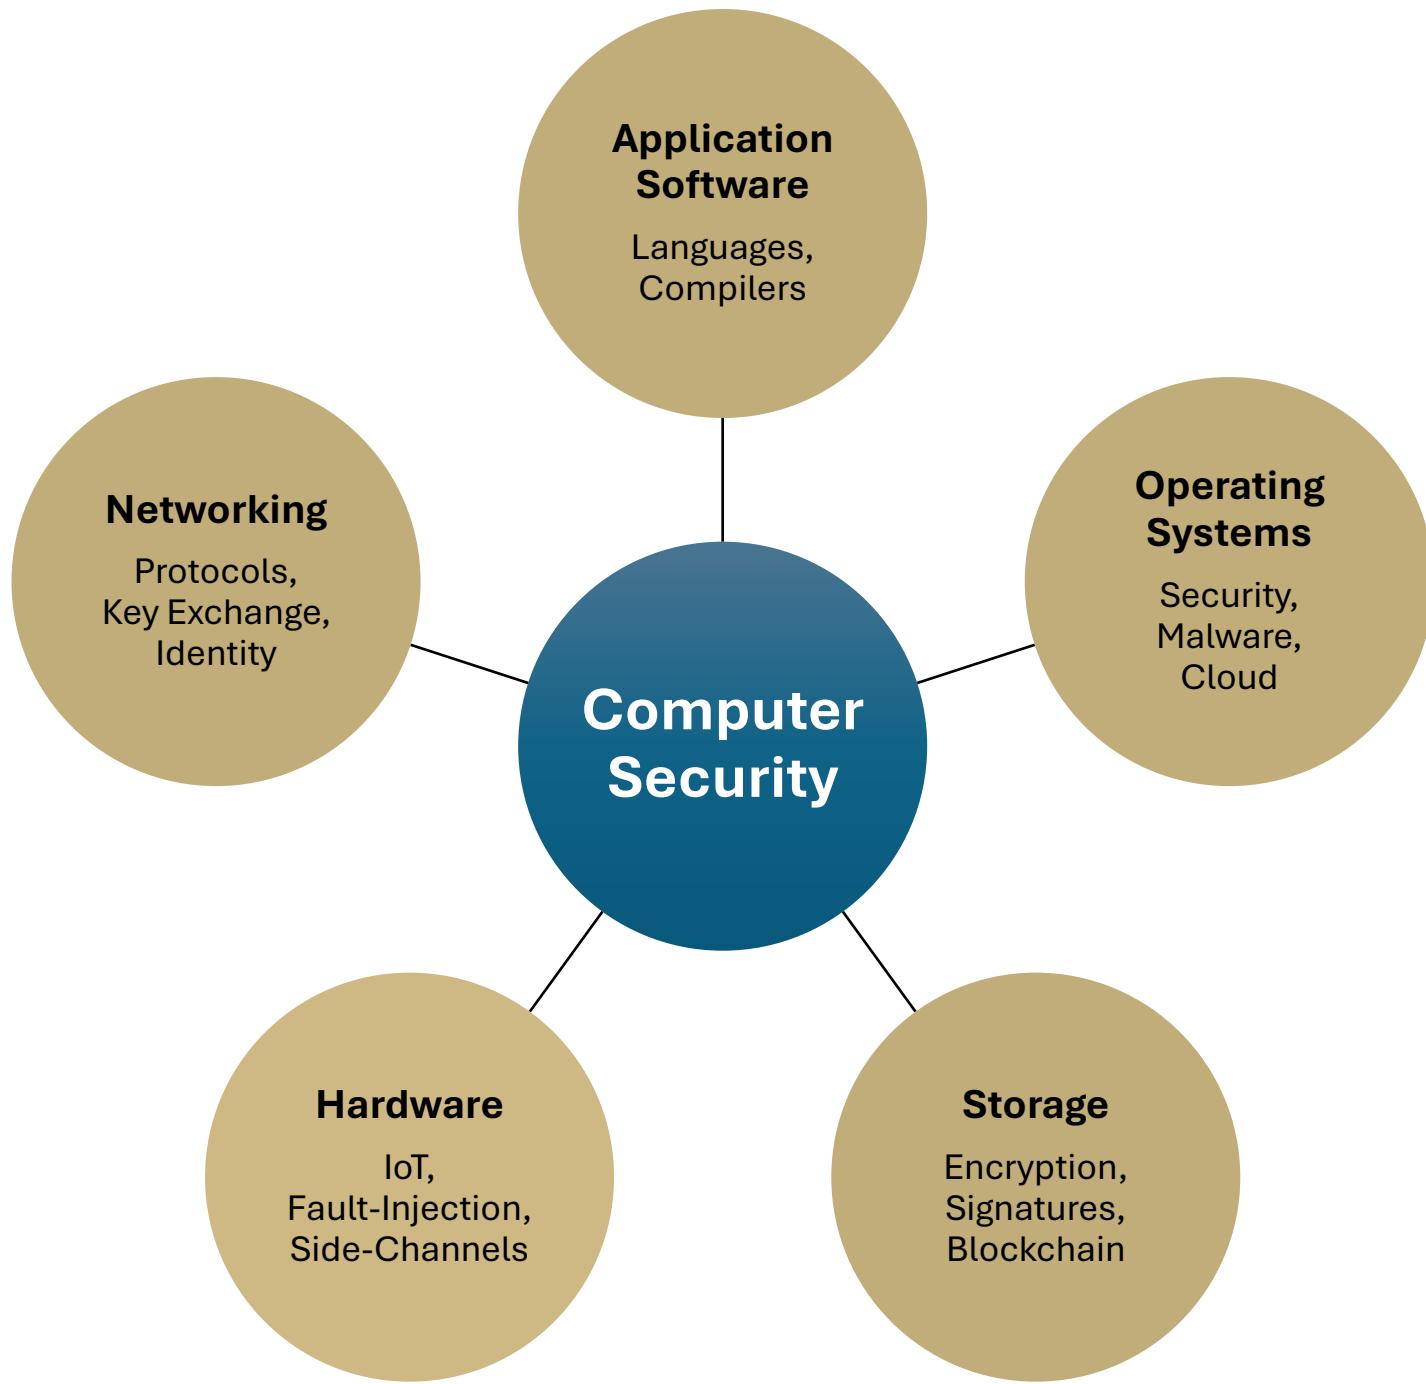
- Law #1: If a bad guy can persuade you to run his program on your computer, it's not solely your computer anymore.
- Law #2: If a bad guy can alter the operating system on your computer, it's not your computer anymore.
- Law #3: If a bad guy has unrestricted physical access to your computer, it's not your computer anymore.
- Law #4: If you allow a bad guy to run active content in your website, it's not your website any more.
- Law #5: Weak passwords trump strong security.
- Law #6: A computer is only as secure as the administrator is trustworthy.
- Law #7: Encrypted data is only as secure as its decryption key.
- Law #8: An out-of-date antimalware scanner is only marginally better than no scanner at all.
- Law #9: Absolute anonymity isn't practically achievable, online or offline.
- Law #10: Technology is not a panacea.

Law #1: If a bad guy can persuade you to run his program on your computer, it's not solely your computer anymore

It's an unfortunate fact of computer science: when a computer program runs, it will do what it's programmed to do, even if it's programmed to be harmful. When you choose to run a program, you are making a decision to turn over a certain level of control of your computer to it - often anything up to the limits of what you yourself can do on the computer (and sometimes beyond). It could monitor your keystrokes and send them to criminals eager for the information. It could open every document on the computer, and change the word "will" to "won't" in all of them. It could send rude emails to all your friends. It could install a virus. It could create a "back door" that lets someone remotely control your computer. It could relay a bad guy's attack on someone else's computers. Or it could just reformat your hard drive.

That's why it's important never to run a program from an untrusted source, and to limit the ability of others to make that decision for you on your computer. There's a nice analogy between running a

Microsoft's Ten Immutable Laws of Security





Course Overview



UNIVERSITY OF
TORONTO

ECE568

- We will be focusing on the technology used to **build software and systems**.
- We will **not** be focusing on configuration settings, administrative policies or physical controls:
 - Wifi security standards
 - Data retention policies, data-destruction procedures
 - Cooling / fire-suppression, guards, fences, cameras, etc..
- We will look at cryptography from the standpoint of how to build secure systems (rather than focusing on formal proofs).
- You have a responsibility to use what we learn in this course to **create better security, rather than break it**.

Quercus

<https://q.utoronto.ca/>



Syllabus
Lecture Slides
Assignments
Past Tests
Announcements

Course Outline

Part 1: Course Overview and Review

- Course introduction
- Quick review: O/S, assembly language, software tools
- Basic principles of computer security
- Confidentiality, Integrity, Availability

Primary Learning Goals:

1. Ability to explain the fundamental goals of Confidentiality, Integrity and Availability, and how they each can be applied to creating secure product designs.
2. Ability to use basic software debugging tools (e.g., gdb) to analyze program behaviour.

Course Outline

Part 2: Software Code Vulnerabilities

- Injection-type attacks: buffer overflows, format strings, double-free, ROP, etc.
- Library attacks: encapsulation, etc.
- Defenses and good programming practice

Primary Learning Goals:

1. Ability to explain and execute the most common attacks against software vulnerabilities, in order to locate design and implementation flaws.
2. Ability to use appropriate tools and best-practice coding techniques to protect against these types of attacks in your own work.

Course Outline

Part 3: Cryptography

- Encryption: Symmetric (block/stream ciphers) and Public Key
- Secure key-exchange
- Hashes and Signatures: MAC, HMAC, blockchain, etc.
- Secure Communication Protocols: SSL/TLS, VPN

Primary Learning Goals:

1. Ability to select appropriate encryption techniques to meet both the Data Confidentiality requirements and the performance requirements of your software.
2. Ability to securely exchange cryptographic keys over insecure network channels.
3. Ability to use hash-based tools to provide Data Integrity and Data Authentication, through the application of MACs, HMACs and digital signatures.

Course Outline

Part 4: Web Security

- Web authentication, cookies
- Web attacks: Cross-Site Scripting (XSS), CSRF, etc.
- SQL Injection
- Cloud security

Primary Learning Goals:

1. Ability to safely use web cookies to secure sessions and end-user data.
2. Ability to execute the most common web-based attacks, in order to locate design and implementation flaws.
3. Ability to use appropriate tools and best-practice coding techniques to protect against these types of attacks in your own work.

Course Outline

Part 5: Access Control

- Access control models
- Hardware Security Modules (HSMs)
- Side-Channels and Covert Channels
- Attacks on IoT/embedded systems

Primary Learning Goals:

1. Ability to identify potential sources of side- and covert-channels that could be used to leak information out of both hardware- and software-based systems.
2. Ability to identify and describe potential injection vulnerabilities, and how they might compromise the security of IoT/embedded systems.

Course Outline

Part 6: Network Security

- Network-layer security risks: BGP, ARP, DNS, TCP/IP, etc.

Primary Learning Goals:

1. Ability to identify how vulnerabilities in many standard networking protocols could be exploited by a third-party to pose risks to confidentiality, integrity and availability of your software systems.
2. Ability to use appropriate tools and best-practice techniques to protect your own systems and designs.

Course Outline

Part 7: Malware

- Viruses and worms
- Rootkits and bootkits

Primary Learning Goals:

1. Ability to identify the mechanisms through which computer viruses, worms, rootkits and bootkits functions and propagate, and the risks these may pose to your systems.

Course Outline

Section 8: Physical Security

- Physical security system design and vulnerabilities

Primary Learning Goal:

1. Ability to explain the most common vulnerabilities and defenses in physical security systems, and relate this to design challenges in digital security designs.

Optional Textbooks

Security in Computing

- Pfleeger and Pfleeger

Computer Security: Principles & Practice

- Stallings and Brown

Applied Cryptography

- Bruce Schneier

Marking Scheme

- **Labs:** 25%
- **Midterm:** 35%
- **Final Exam:** 40%

The midterm and final are both “Type C”
(single reference sheet), no calculator.



Plagiarism

All labs, assignments and tests are to be completed with your own original work.

Anything submitted for credit must be something that you (and your lab partner) produced.



What to Expect

Course Covers a Lot of Material

- Memory, CPUs
- Operating Systems, Networking
- Cryptography

Where You'll Spend Your Time

- Four lab assignments
- Most of the work will be in the labs: course focuses on practical aspects of security

Lab Assignments

Lab 1: Identification of Vulnerabilities, Construction of Attacks

- You will be given some vulnerable programs.
- Your job is to construct attacks that will let you hijack the programs and spawn a command shell.

Lab 2: Cryptography

- You will use X509 certificates and digital signatures to create a basic blockchain.

Lab 3: Web Security

- You will be given a web application with a number of security vulnerabilities.
- You will craft a series of attacks that exploit the vulnerabilities in the application's design.

Lab 4: Network Security

- TBD (Under Revision)



Computer Security

Rules and Assumptions



UNIVERSITY OF
TORONTO

Rules

Like the programs that enforce them, computer systems have rules.

Some rules are **explicit** and well thought out; others are **implicit** and based on assumptions.

Implicit rules are often a source of security vulnerabilities.



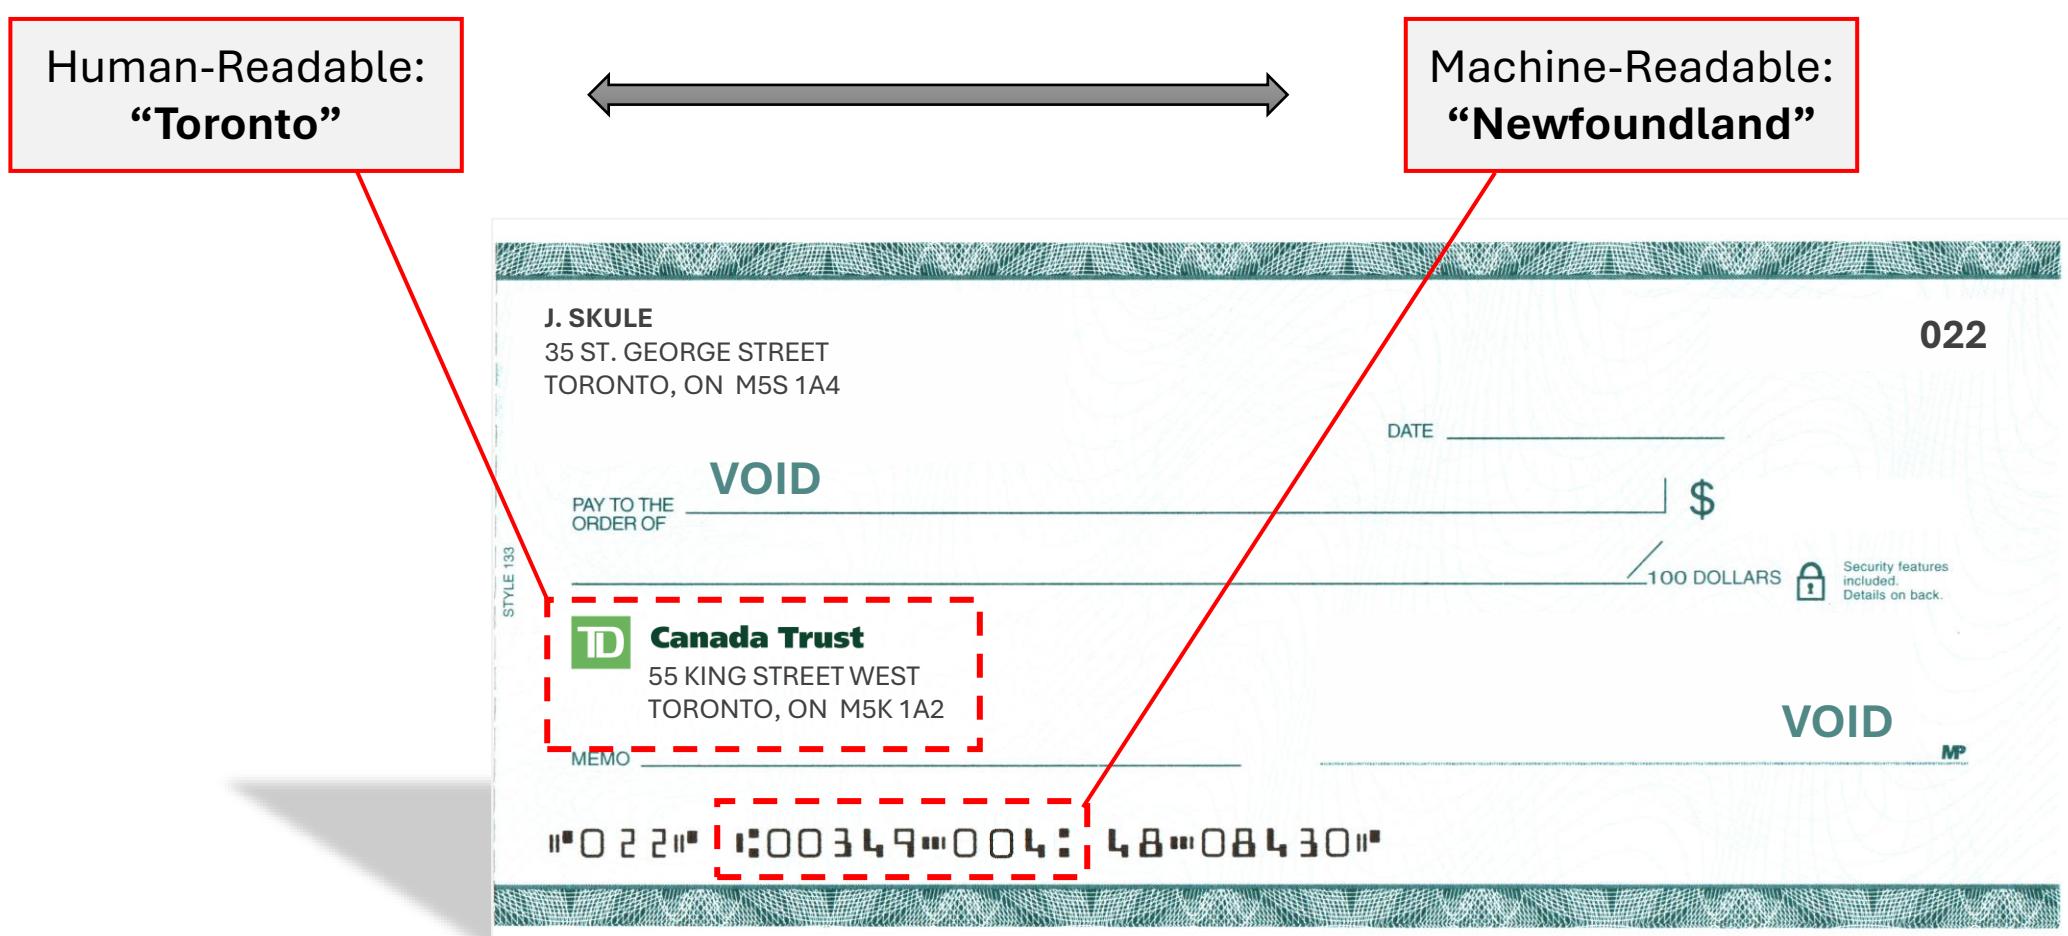
Assumptions

Security risks occur when our assumptions turn out to be false:

- Data
- Input
- User behaviour



Case Study: “Cheque Kiting” Scam (1980’s)





Computer Security

What is it? What makes it difficult?



UNIVERSITY OF
TORONTO

Computer Security

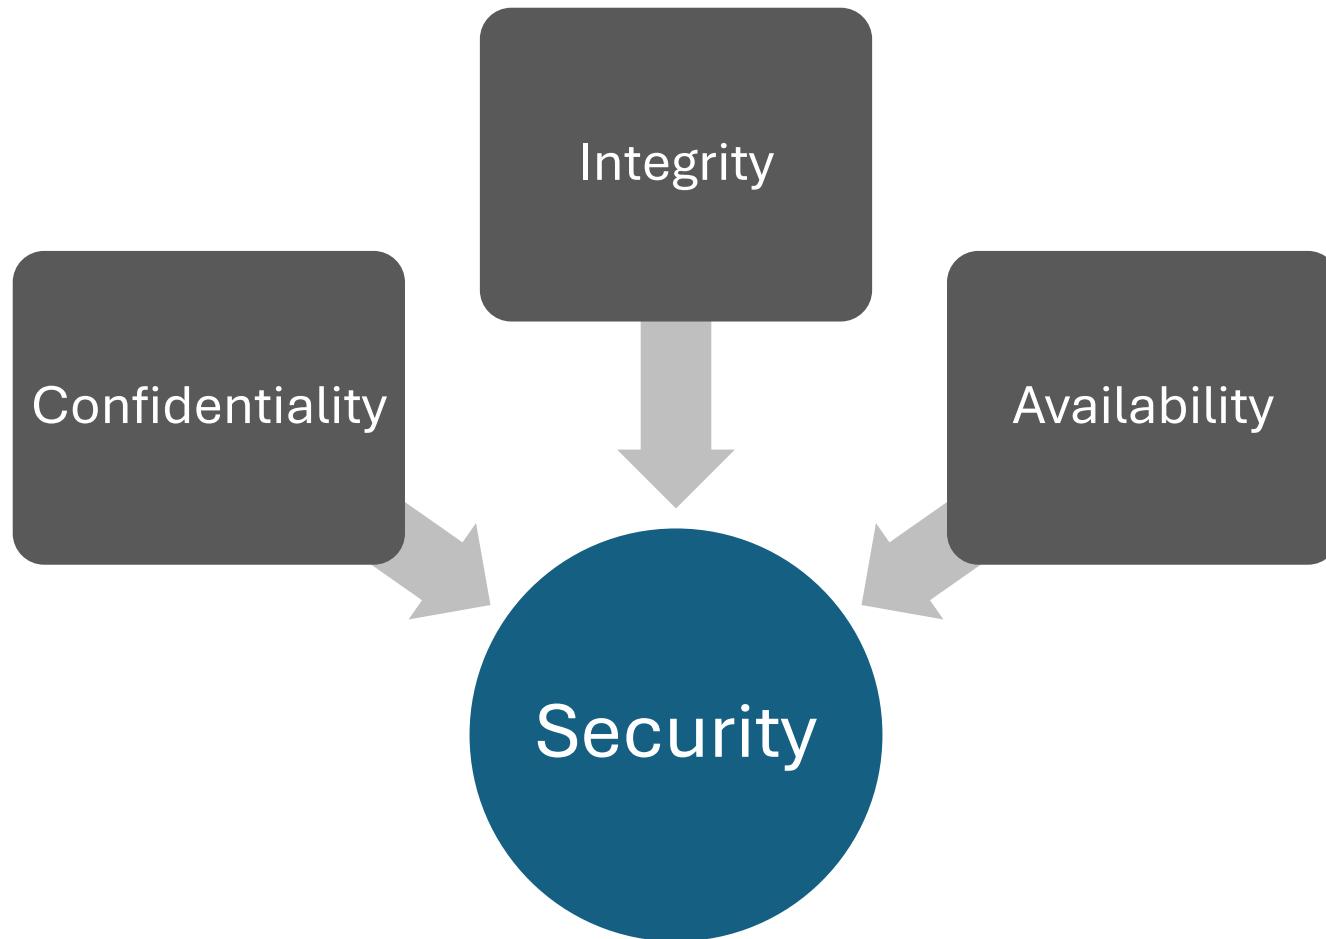
Cybersecurity is about **understanding** a system really well, and **questioning** the implicit rules.

- A **reliable** system does what it is supposed to do.
- A **secure** system does what it is supposed to do,
and nothing else.

Cybersecurity is...

1. ...the **protection** of information and information-systems...
2. ...against **unauthorized** access, use, disclosure, modification or destruction...
3. ...in order to provide **confidentiality, integrity** and **availability**.

Security Fundamentals



Confidentiality

Confidentiality is the protection of data or resources from exposure.

There are two aspects of information or resources that are often important to conceal:

1. The **content** of the data or resource (e.g., passwords, etc.).
2. The **existence** of the data or resource (as merely knowing that something exists may be damaging).

- **Organizational Controls**

- Acceptable-Use Policies
- **Security Training**
- Privacy Policies
- Retention Policies

- **Access Rules**

- Software
- **Hardware**

- **Cryptography**

- Authentication / authorization
- **Encryption**

Integrity

Integrity is establishing and maintaining the trustworthiness of the data or resource.

There are two important aspects of a resource or piece of data:

1. Correctness of its **contents**: can we tell if it's been altered?
2. Correctness of its **origin**: can we authenticate who this really came from?

- **Tools**
 - Monitoring (anti-malware, etc.)
 - Backups
- **Access Rules**
 - Policies
 - Controls (SQL Injection, etc.)
- **Cryptography**
 - Hashing
 - Signing
 - Cloud-specific tools/techniques

Availability

Availability ensures the data or resource can be used as/when desired.

- A **resource** is available if it is responding to requests
- **Data** are available if the service that stores the information is up and running

Availability is generally harder to ensure.

- A resource could be temporarily unavailable, but at what time does it really become unavailable?

- **Disaster Recovery**

- Files (data/software)
- Servers
- Networking

- **Backups**

- **Mitigation Tools**

- Firewalls
- Denial-of-Service prevention

Security: CIA

- All secure systems provide **Confidentiality, Integrity and Availability** to some degree
- However, no system is absolutely secure, and no system can provide all (any?) of the three absolutely
- Generally, a system is considered secure if the security requirements are met **for a given time**
 - e.g., confidential data is not leaked until it is declassified

Vulnerabilities

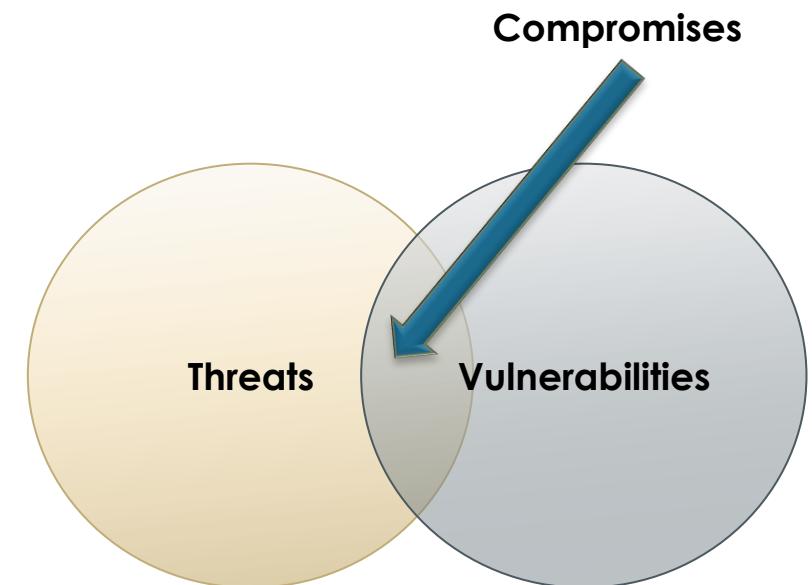
A **vulnerability** is a flaw that weakens the security of the system. Vulnerabilities are very difficult to identify even after significant amount of testing

- They can be serious:
 - An unchecked string copy allows an attacker to overflow a buffer and execute arbitrary code in a privileged program
 - During configuration, a system administrator forgets to disable debug mode on a program, allowing an attacker to gain privileged information
 - A naïve user does not change the default password on their router from the factory default

Threats

A **threat** is any method that can be used to potentially breach security and cause harm.

- When an attacker exercises a threat, it becomes an **attack** (also called an “exploit”)
- If the attack is successful, the system is **compromised**



Security is a Negative Goal

Ensuring that something **cannot** happen is much harder to design, measure and verify than traditional software quality-assurance goals.

- **Positive goal:** Alice can read the file
- **Negative goal:** Bob cannot read the file

Problem: In what ways might an attacker try to access Alice's file? (Not an easy question to answer!) Start with designing **rules**, and then check the assumptions those rules rely on.



Review

Memory: Data Representation



UNIVERSITY OF
TORONTO

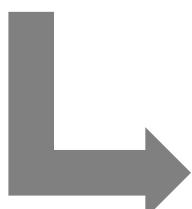
Data Representation

As developers, our primary job
is to manipulate numbers.



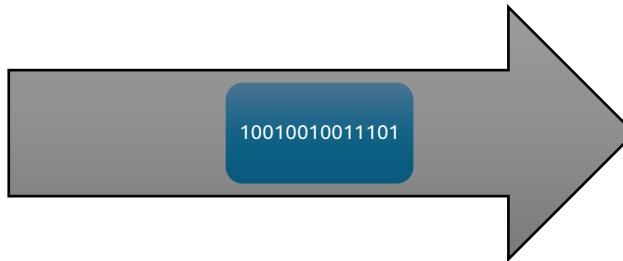
DEMO

```
size_t  
void *  
int  
void *  
pthread_t  
pthread_attr_t  
int  
struct main_args  
int *  
  
stackSize = STACK_SIZE;  
targetStack = (void *)STACK_LOCATION;  
pageSize = getpagesize();  
mmap_result = NULL;  
lab_main_thread_id;  
pthread_attr;  
rc = 0;  
args = { argc, argv };  
lab_main_return = NULL;  
  
// Check that our target stack size is an integer multiple of 1  
stackSize += pageSize - ( stackSize % pageSize );  
  
// Allocate a new stack at a fixed location  
mmap_result = mmap(targetStack, ( stackSize + pageSize ),  
PROT_READ|PROT_WRITE|PROT_EXEC, MAP_ANON|MAP_FIXED );  
assert( mmap_result != MAP_FAILED );  
  
// Create a new pthread to run lab_main()  
rc = pthread_attr_init( &pthread_attr );  
assert( rc == 0 );  
rc = pthread_attr_setstack( &pthread_attr, targetStack, stacks );  
assert( rc == 0 );
```



The image shows a large memory dump consisting of a grid of hex values. The values are primarily '63' (ASCII 'A') and '6f' (ASCII 'f'), indicating a repeating pattern of characters. Interspersed among these are other hex values such as '00', '2e', '5f', '73', '62', and '6e'. The grid is tilted diagonally across the page.

```
size_t  
void *  
int  
void *  
pthread_t  
pthread_attr_t  
int  
struct main_args  
int *  
  
// Check that our target stack size is an integer multiple of 16  
stackSize += pagesize - ( stackSize % pagesize );  
  
// Allocate a new stack at a fixed location  
mmap_result = mmap( targetStack, ( stackSize + pagesize ),  
                     PROT_READ|PROT_WRITE|PROT_EXEC, MAP_ANON|MAP_FIXED );  
assert ( mmap_result != MAP_FAILED );  
  
// Create a new pthread to run lab_main()  
rc = pthread_attr_init ( &pthread_attr );  
assert ( rc == 0 );  
rc = pthread_attr_setstack ( &pthread_attr, targetStack, stackSize );
```



Data Representation

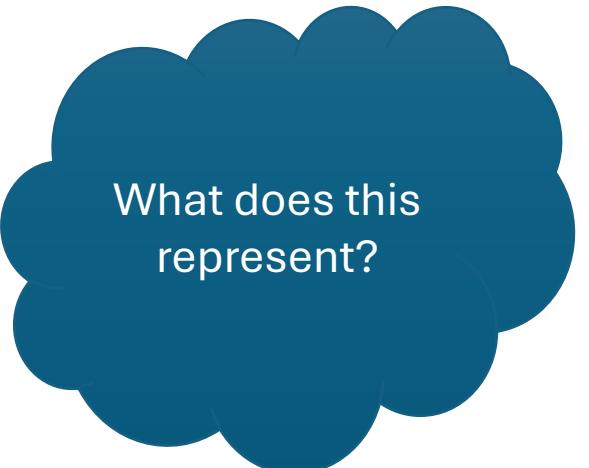
Because everything (instructions, input and output) is represented as numbers, we need to be careful.

Try to avoid thinking of memory only in terms of how you define your variables:

```
int x = 1145258561;
```

0100001010001001000110100100

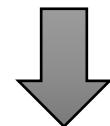
Base 2



What does this
represent?

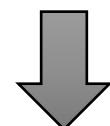
01000001010000100100001101000100

Base 2



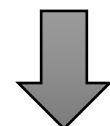
0100	0001	0100	0010	0100	0011	0100	0100
-------------	-------------	-------------	-------------	-------------	-------------	-------------	-------------

Nibbles



0x41424344

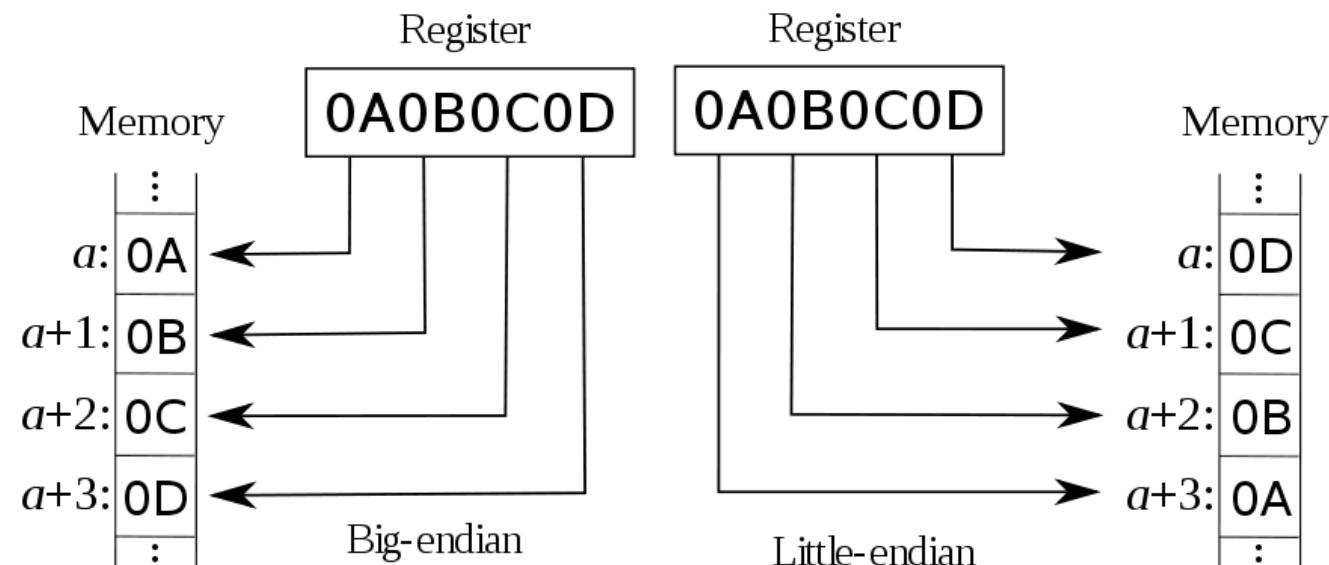
Base 16



1094861636

Base 10

Big-Endian (MSB) vs. Little-Endian (LSB)



Intel is Little-Endian

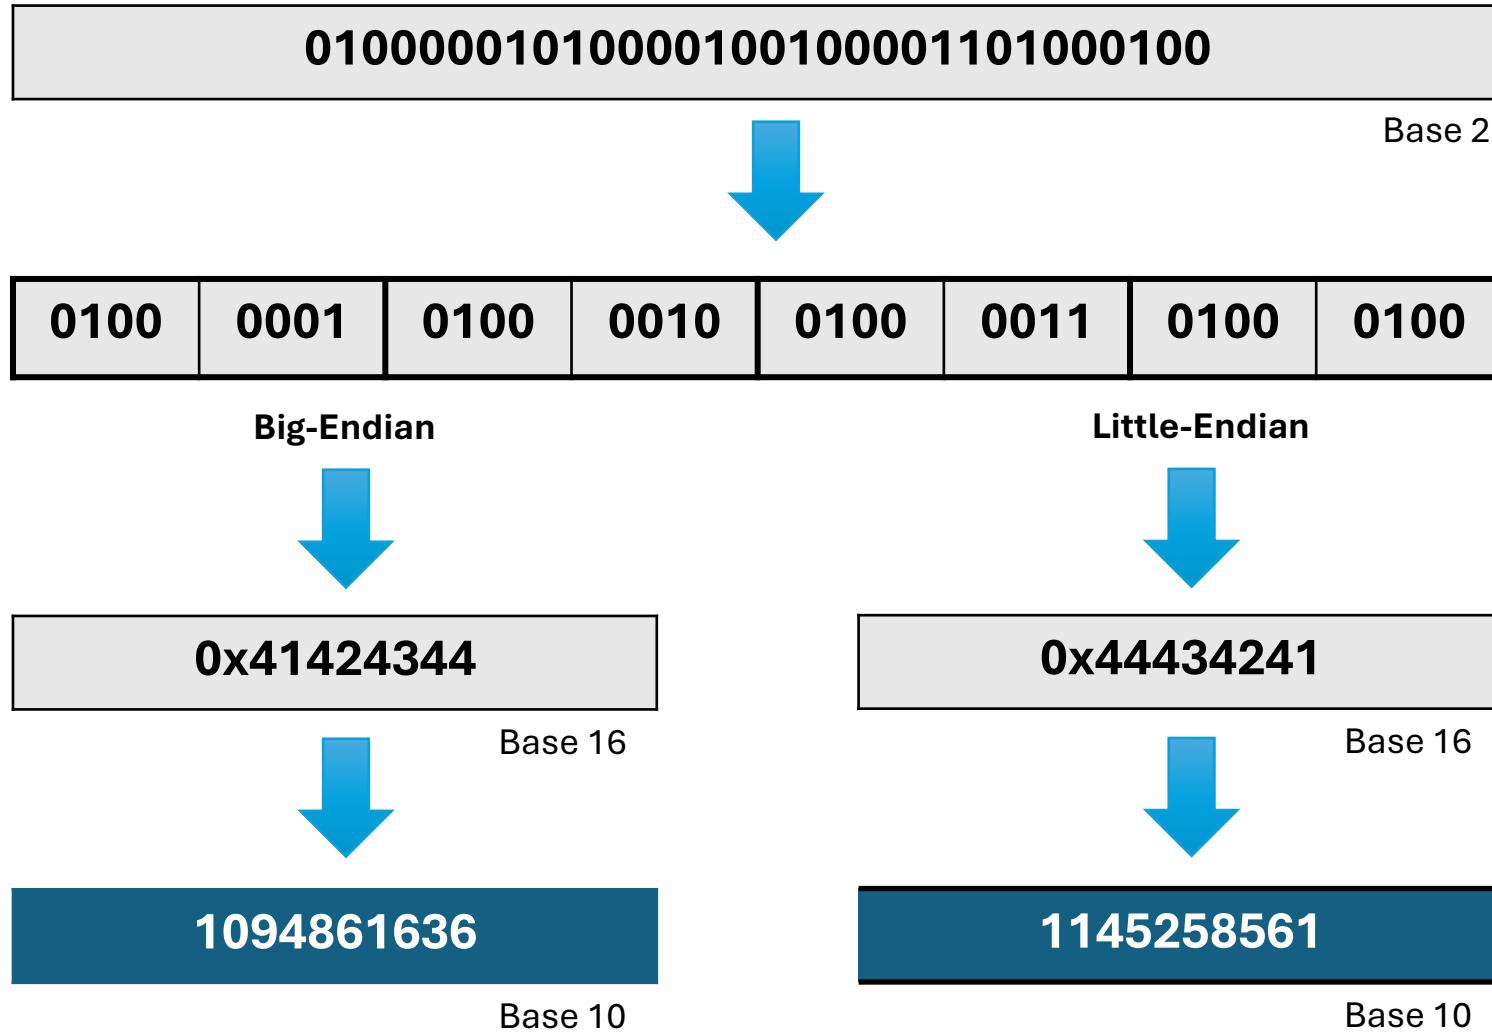
C

```
#include <arpa/inet.h>

int MSB = htonl(0x0A0B0C0D)
int LSB = ntohl(x)
```

Python

```
(0x0A0B0C0D).to_bytes(4,byteorder='big')
```



0100000101000100100001101000100

Base 2



0

1000001

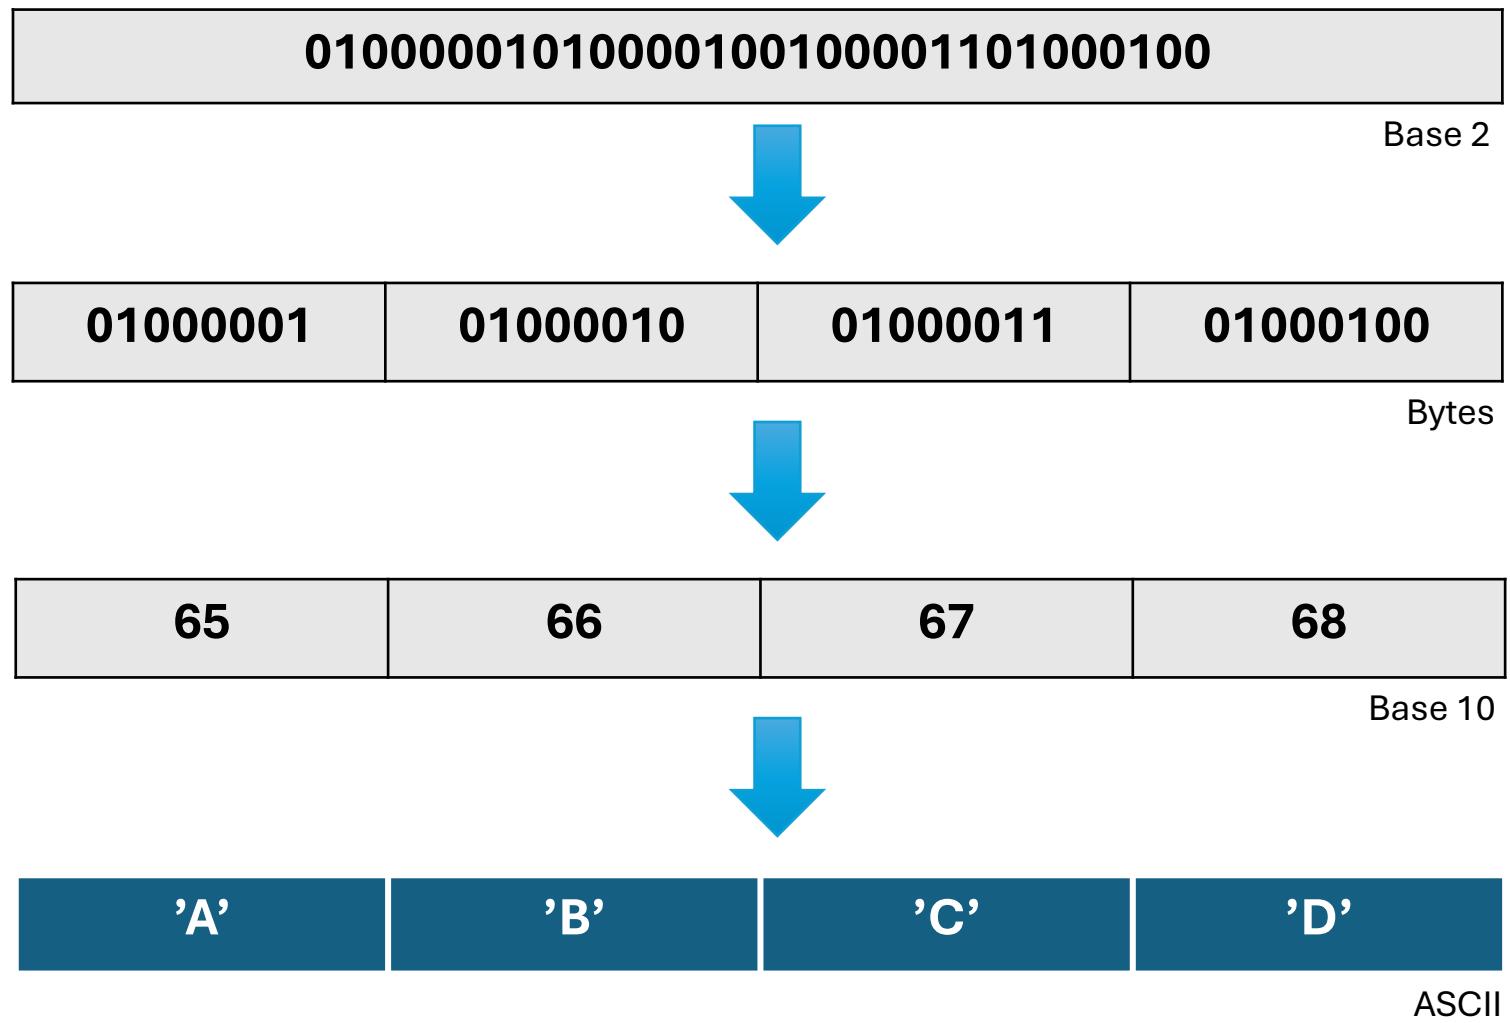
010000100100001101000100

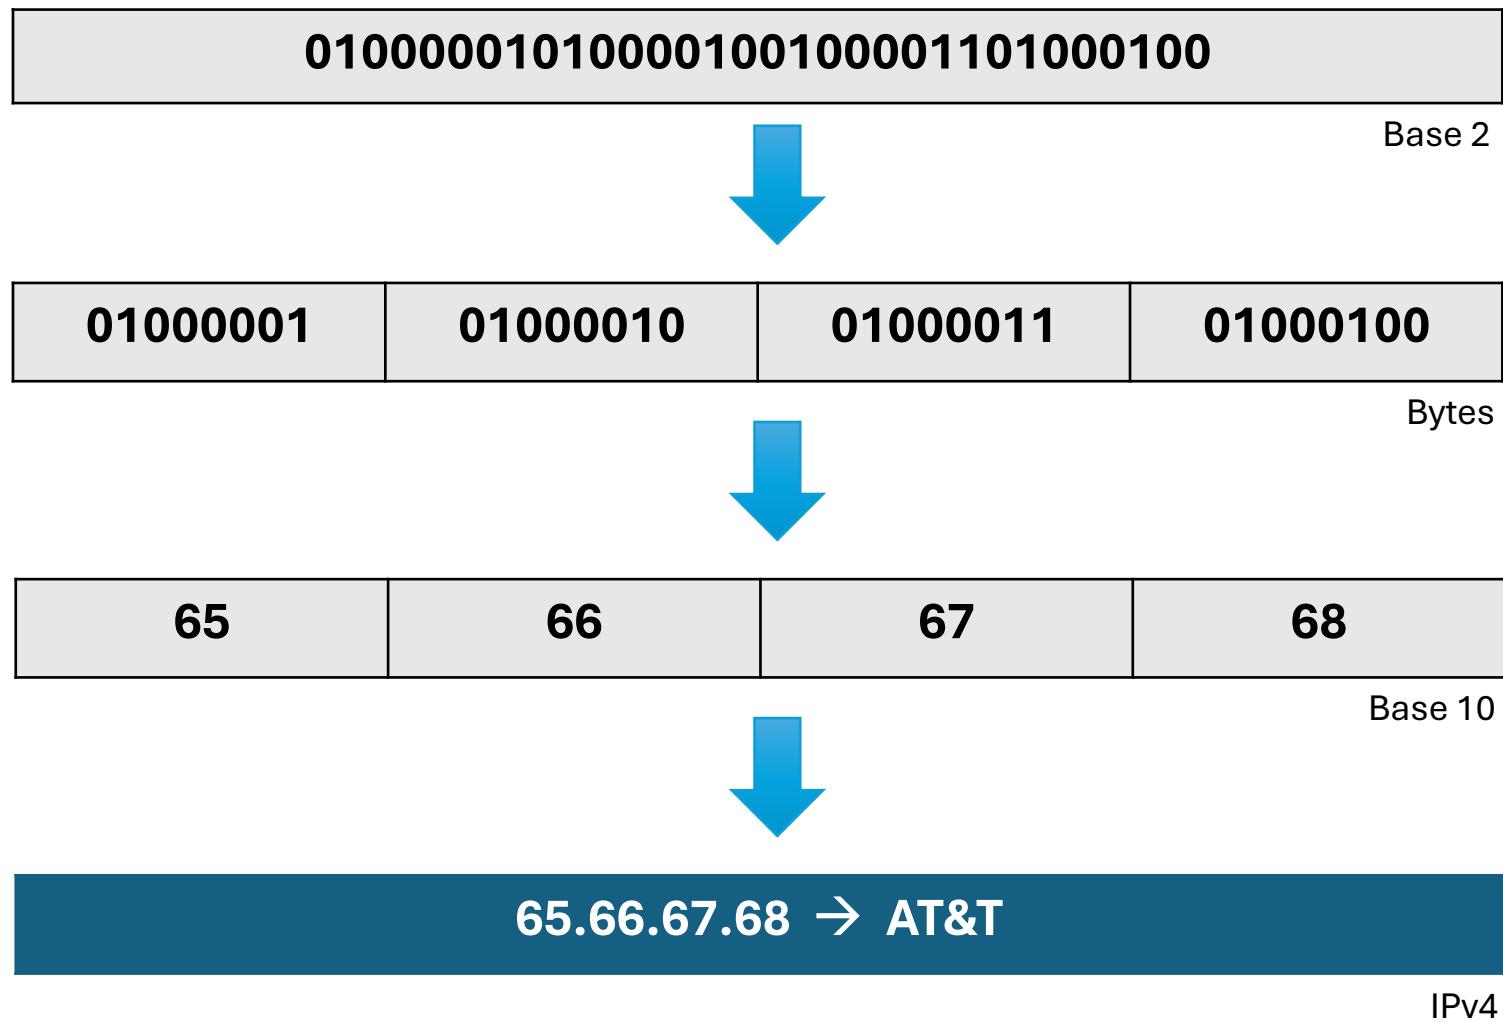
IEEE-754



781.035217

Base 10





0100000101000100100001101000100

Base 2



01000001

01000010

01000011

01000100

Bytes

65

66

67

8-bit Color



Pixel

01000001010000100100001101000100

Base 2



01000001

01000010

01000011

01000100

Bytes



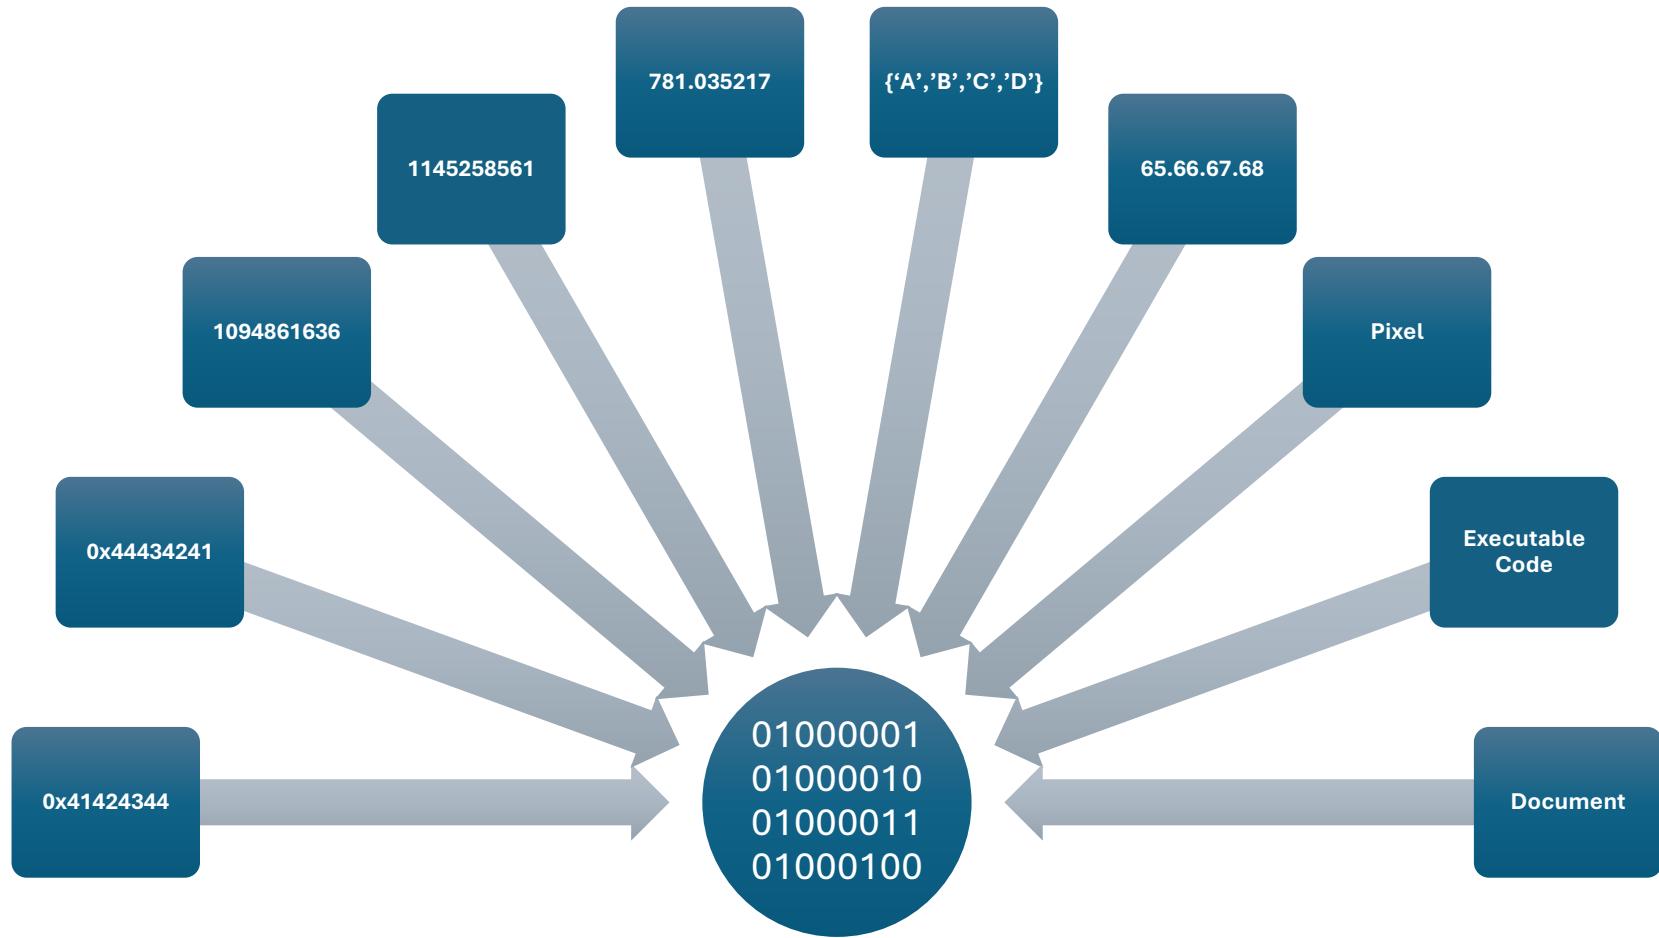
01000001 **rex.B**

01000010 **rex.X**

01000011 **rex.XB**

01000100 **rex.R pop**

Intel Assembly Code

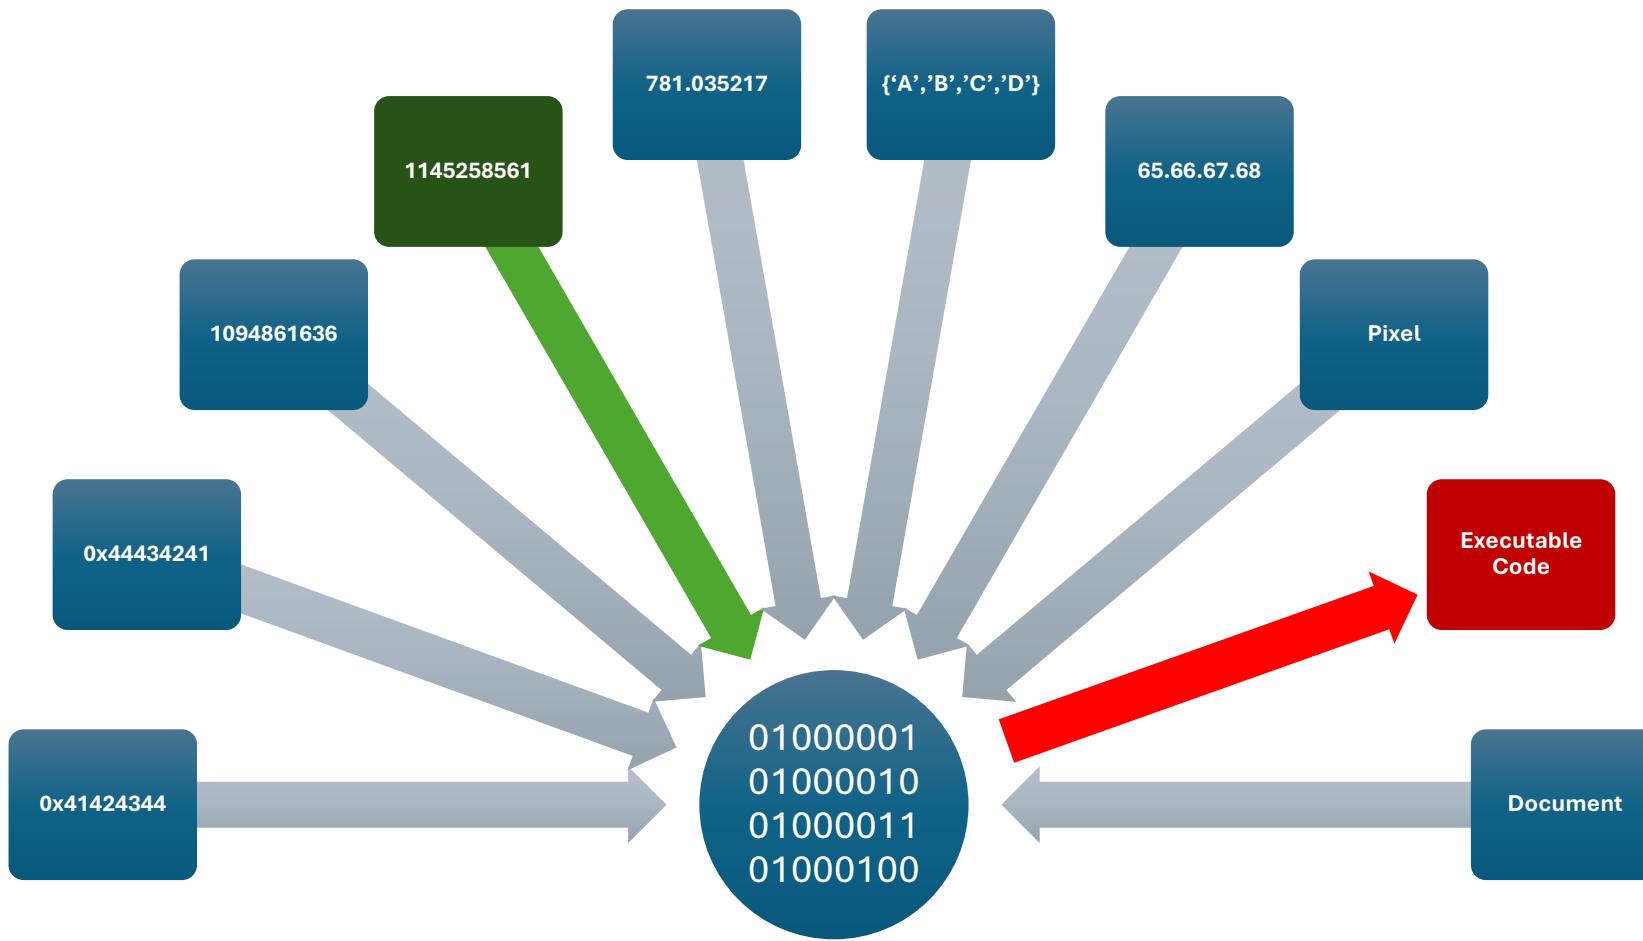


Data Representation

What if our **assumptions** about how the data will be interpreted are **wrong**?



xkcd.com



Many security exploits rely on data being interpreted
in a **different way** than was originally intended.

```
unsigned long int i_hello = 0x6f6c6c6548;  
unsigned long int i_world = 431316168567;  
  
printf("%s %s", (char *)&i_hello, (char *)&i_world);
```

ASCII	Hex	Dec	ASCII	Hex	Dec
A	41	65	a	61	97
B	42	66	b	62	98
C	43	67	c	63	99
D	44	68	d	64	100

Remember: Intel is **Little-Endian**; the bytes in an integer are stored in reverse order in memory.



Review Intel Assembly Code



UNIVERSITY OF
TORONTO

Intel Assembly Code

We will be making use of Intel assembly code, both to discuss and to create various security exploits.

You will not need to be an expert, but should be comfortable reading basic assembly code.

For reference, see:

<http://www.cs.virginia.edu/~evans/cs216/guides/x86.html>

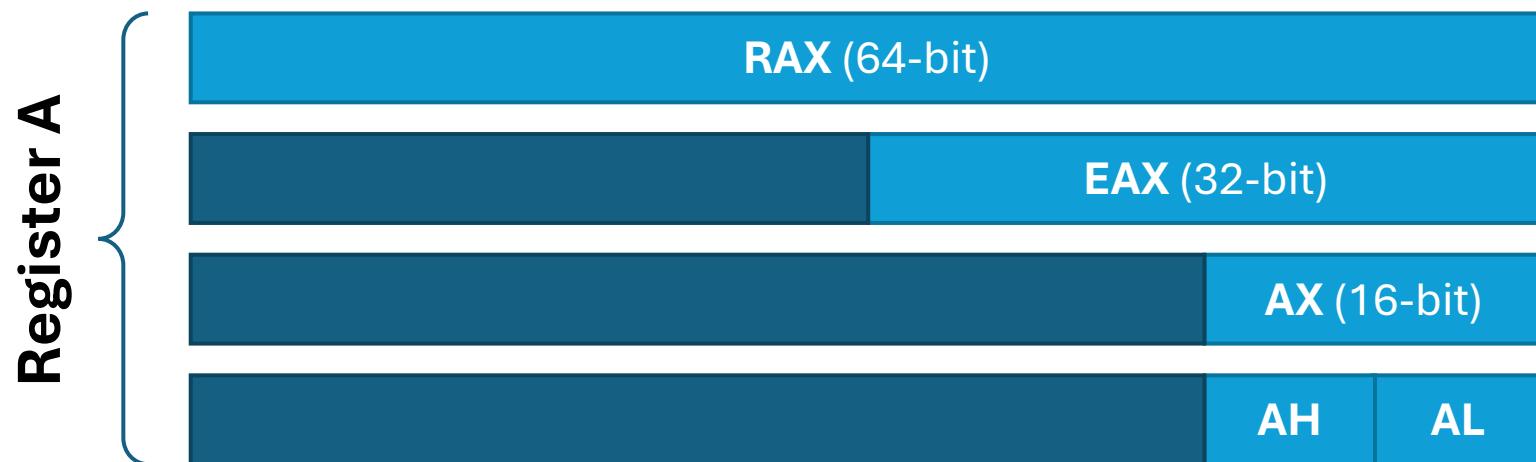
Intel Assembly Code

Intel assembly code looks very similar to most other assembly languages you may already be familiar with (e.g., Motorola 68k)

```
movl -4(%rbp), %eax  
addl $1, %eax  
movl %eax, -16(%rbp)
```

Intel Registers

Current (64-bit) Intel CPUs have six general-purpose registers; the various bits in each individual register can be accessed in a number of ways:



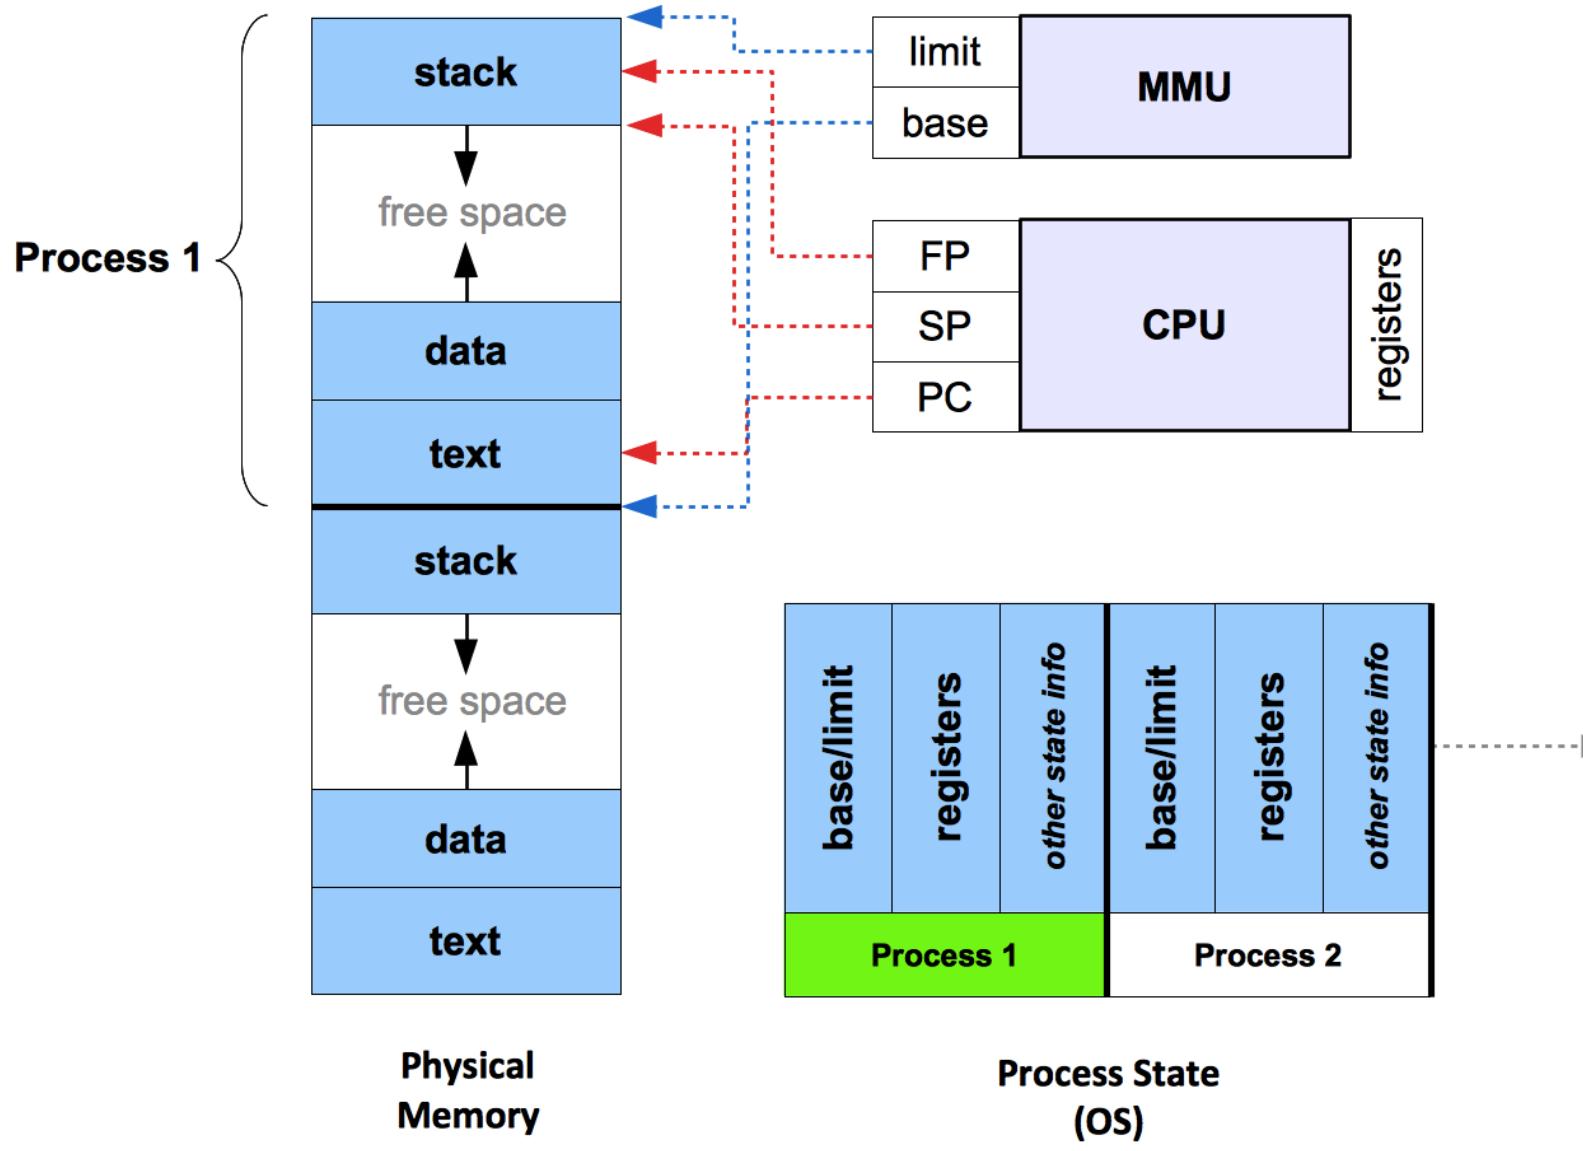


Review

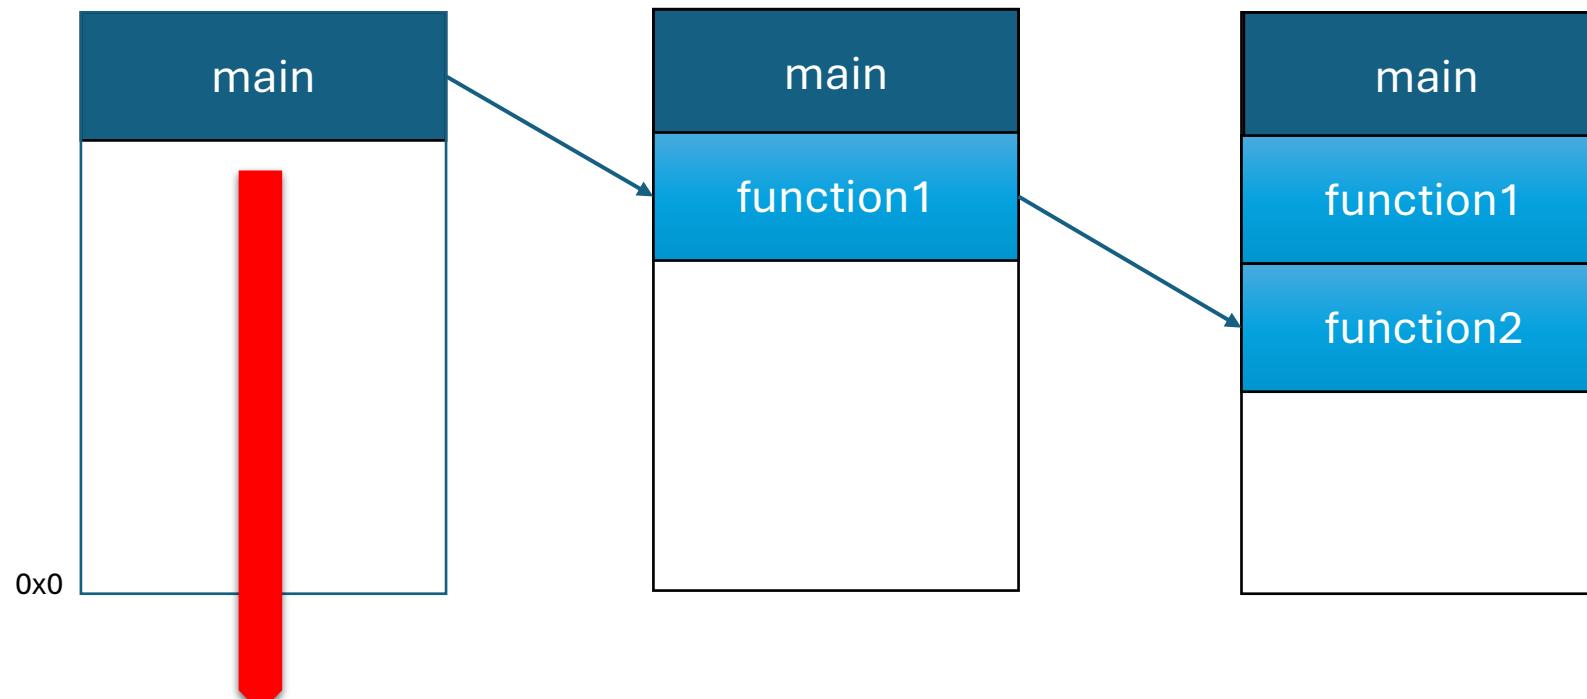
Stack Frames and Function Calls



UNIVERSITY OF
TORONTO

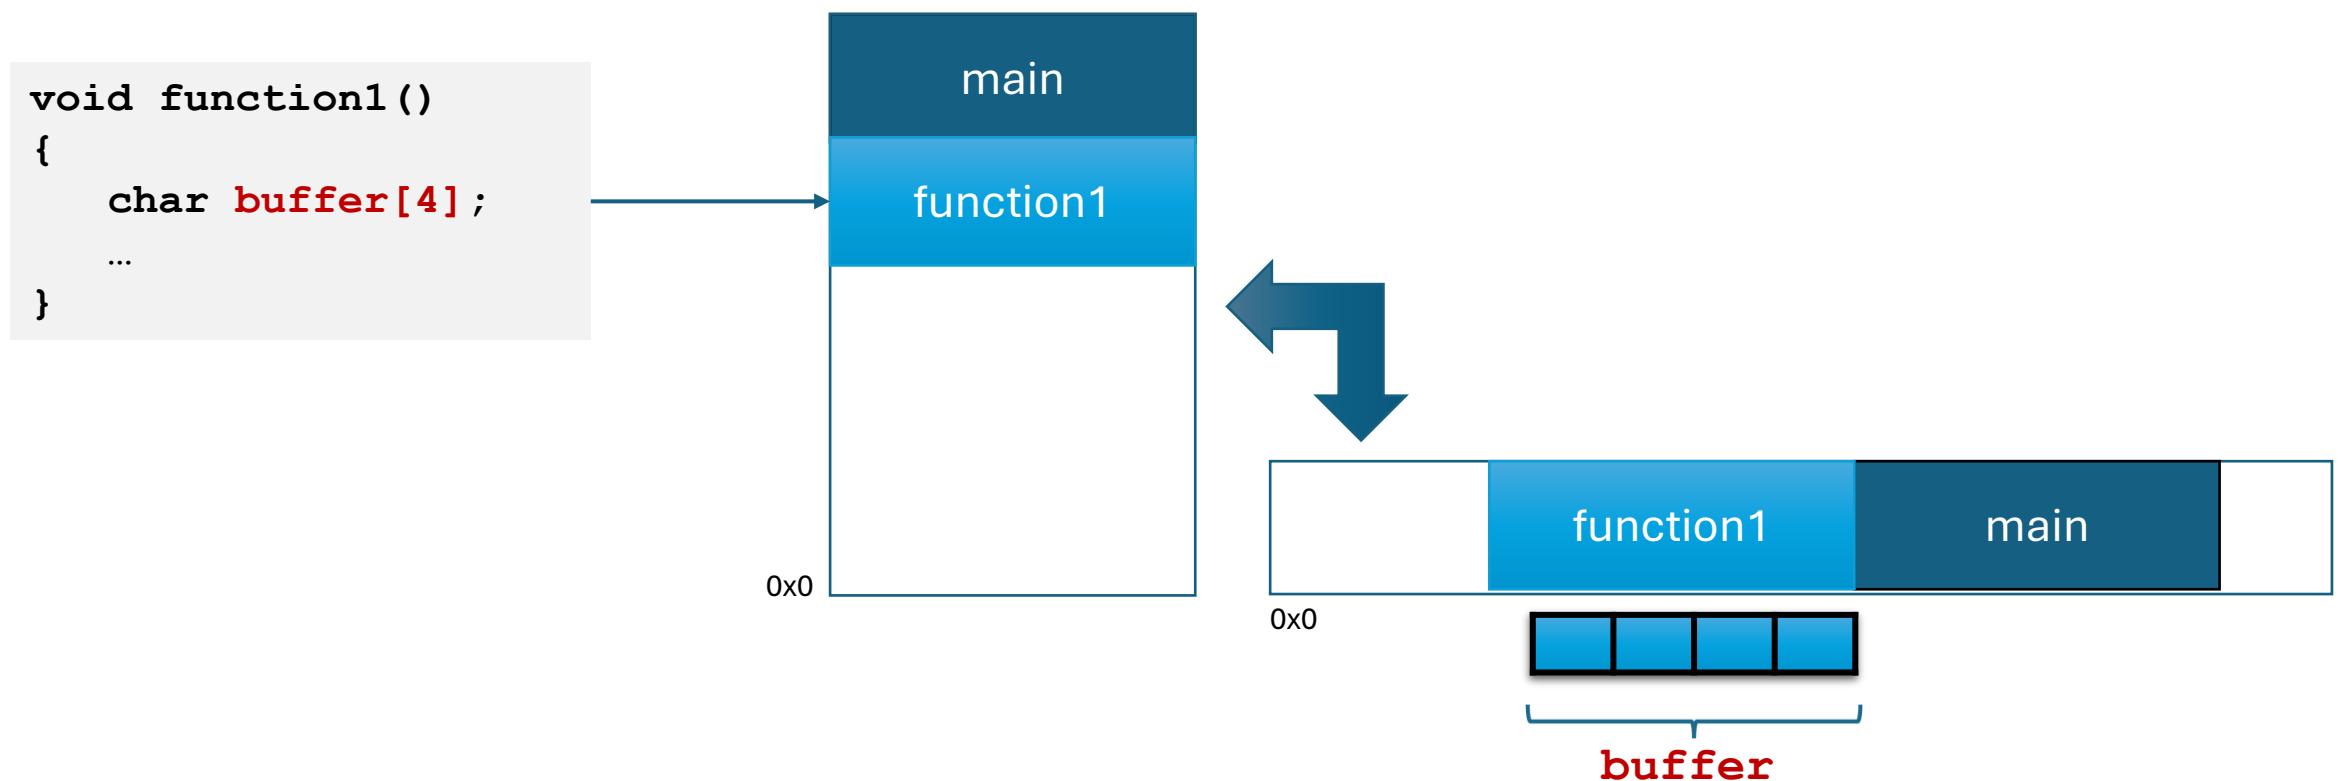


Process Stack

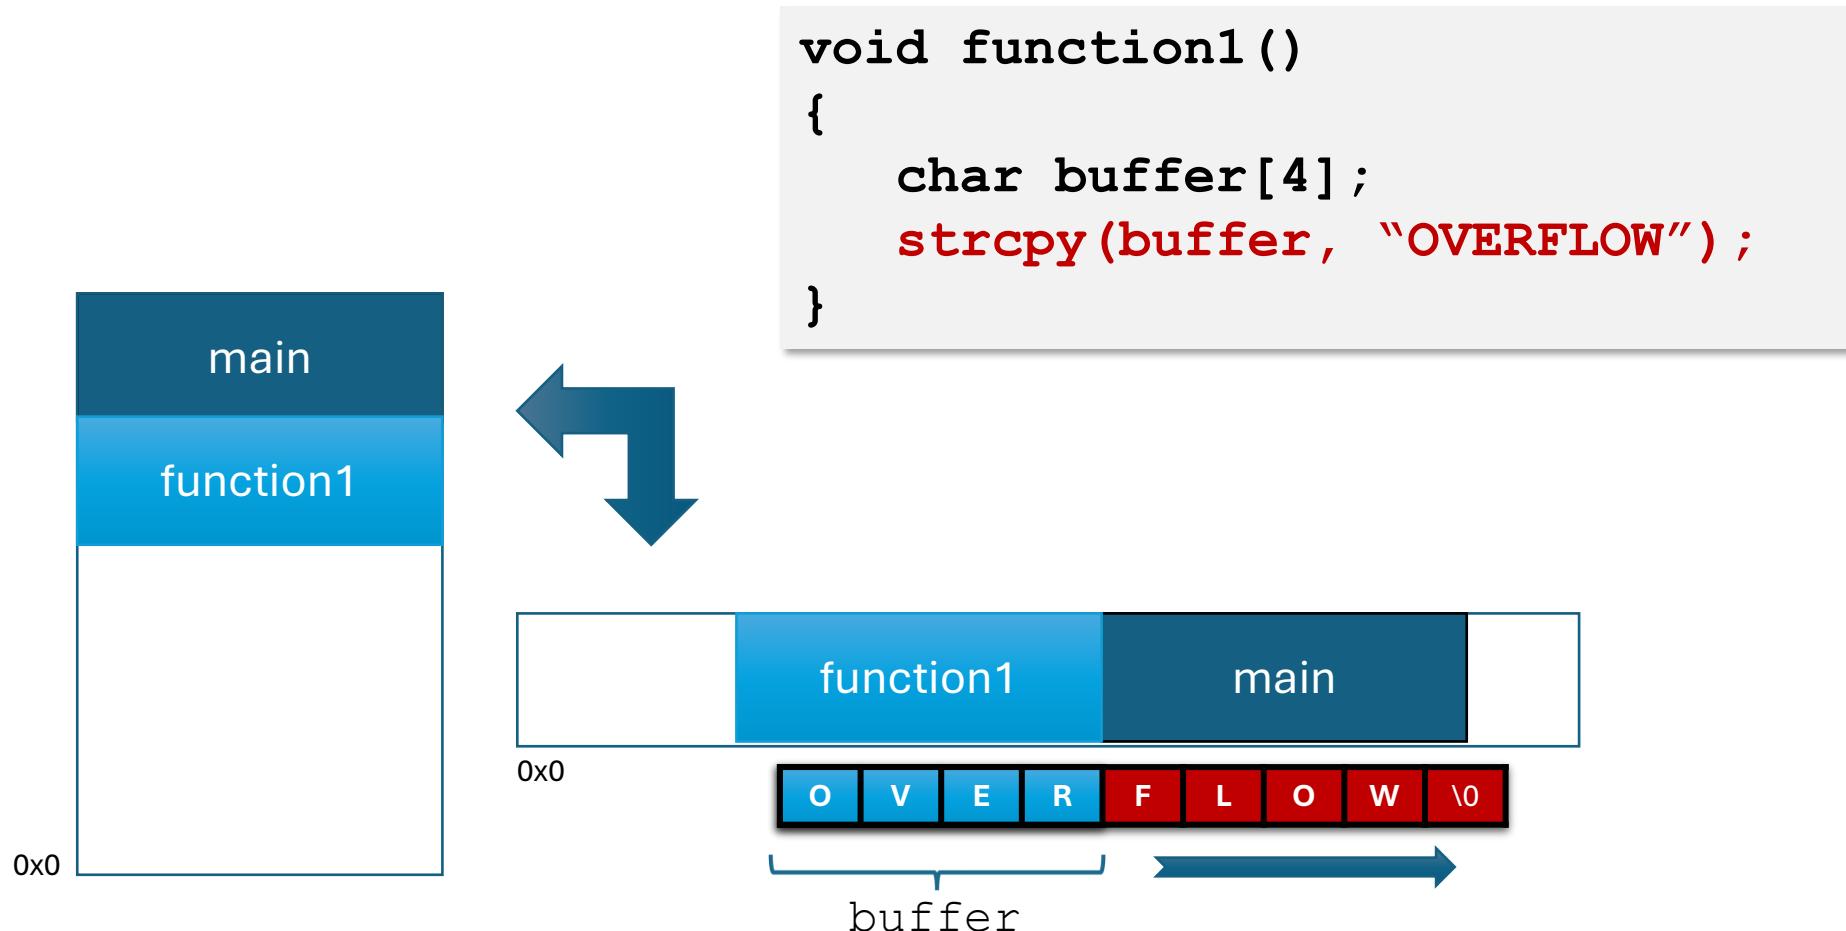


The process stack grows **downwards** on Intel.

Process Stack



Process Stack



Using gdb

break	Defines a new breakpoint
run	Starts a new process
where	List of current stack frames
up / down	Move between frames
info frame	Display info on current frame
info args	List function arguments (“gcc -g”)
info locals	List all local variables (“gcc -g”)
print	Display a variable (“print x”)
x	Display the contents of memory (“x/10x array”)

gdb: info frame

```
(gdb) info frame
```

Stack level 0, frame at 0x7fffffffda20:

rip = 0x401844 in main (test.c:6); saved rip 0x401d14

source language c.

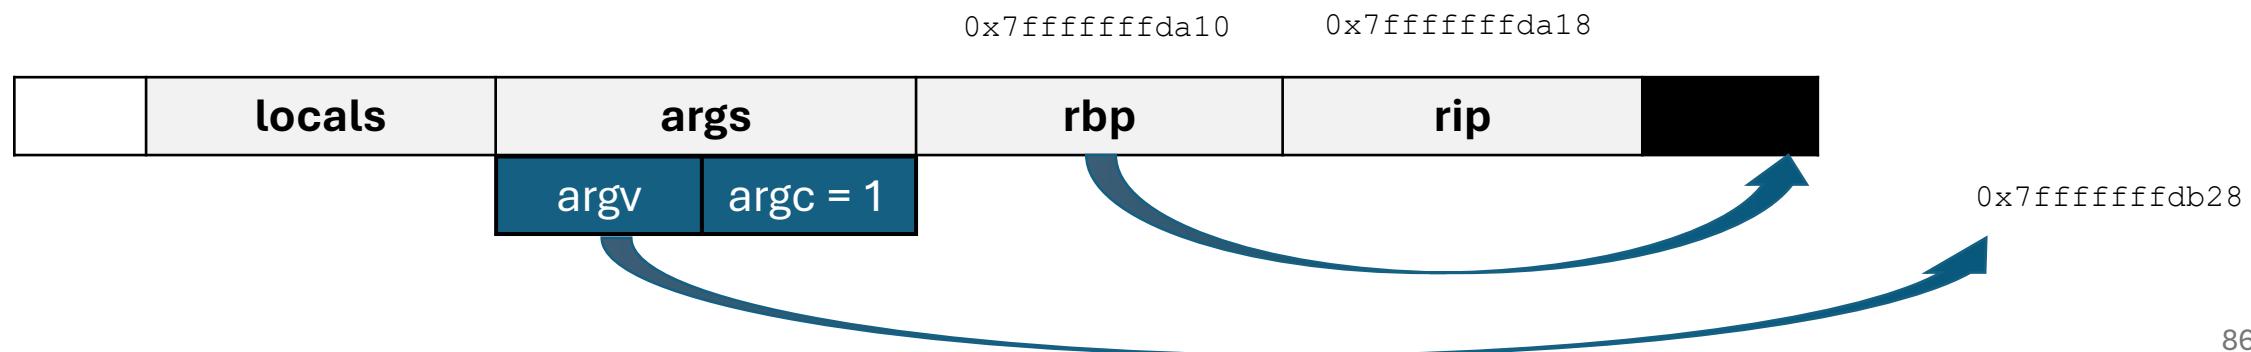
Arglist at 0x7fffffffda10, args: argc=1, argv=0x7fffffffdb28

Locals at 0x7fffffffda10, Previous frame's sp is 0x7fffffffda20

Saved registers:

rbp at 0x7fffffffda10, rip at 0x7fffffffda18

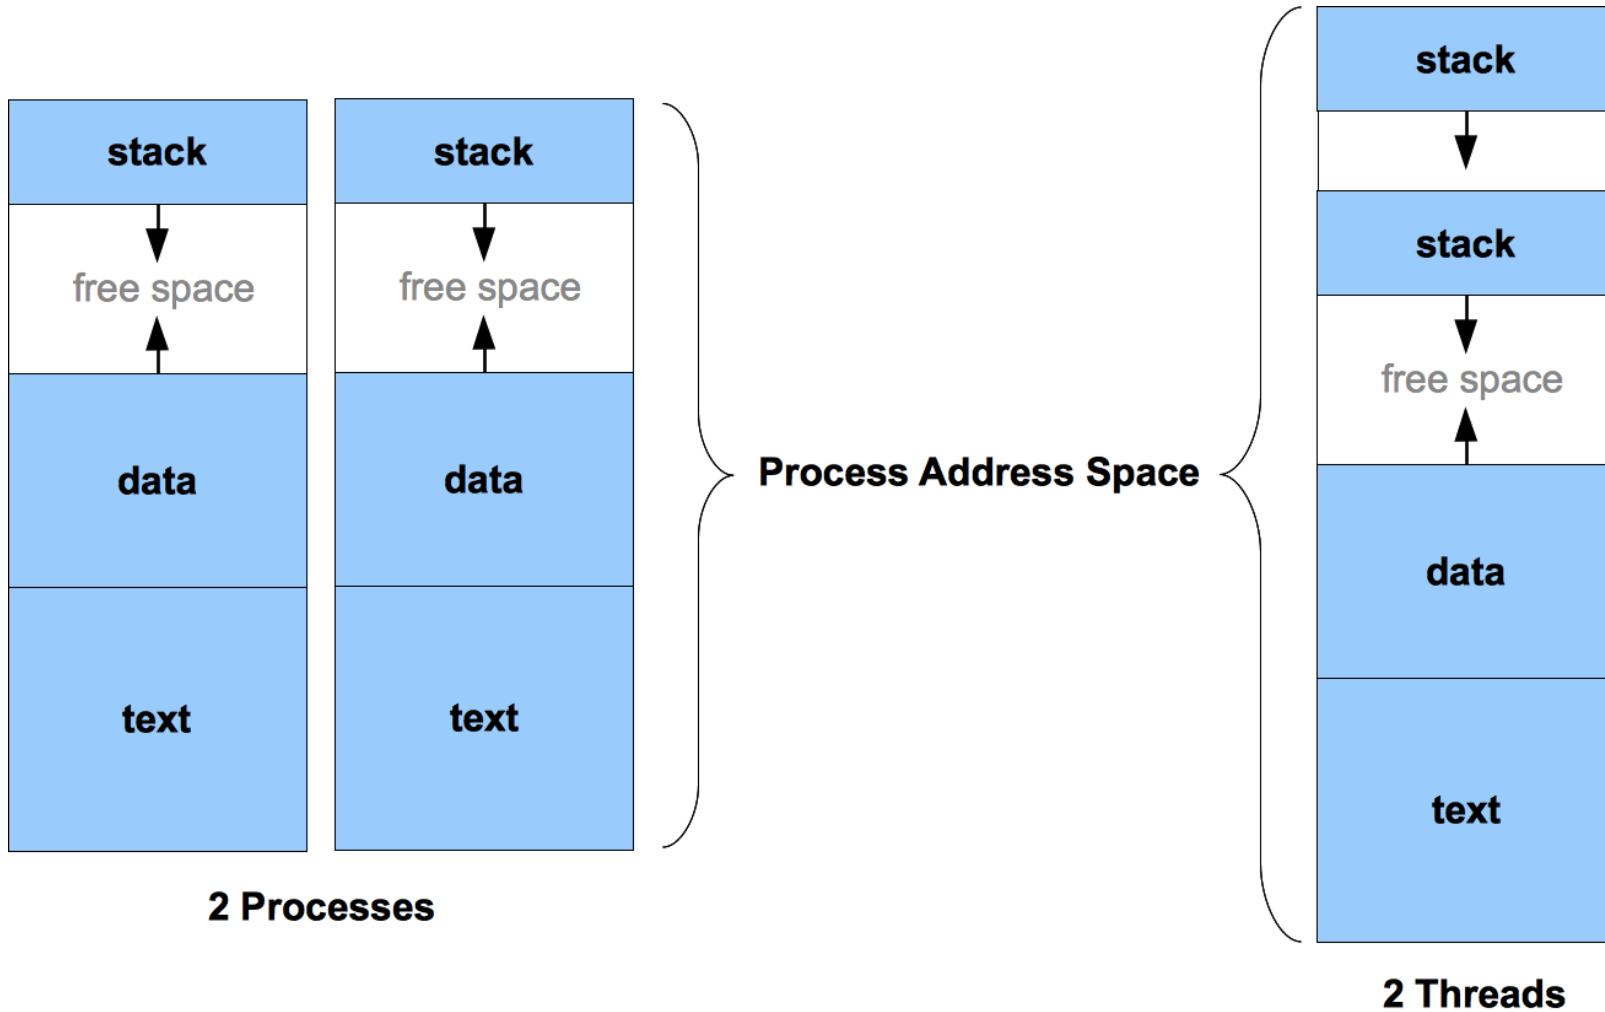
```
(gdb)
```

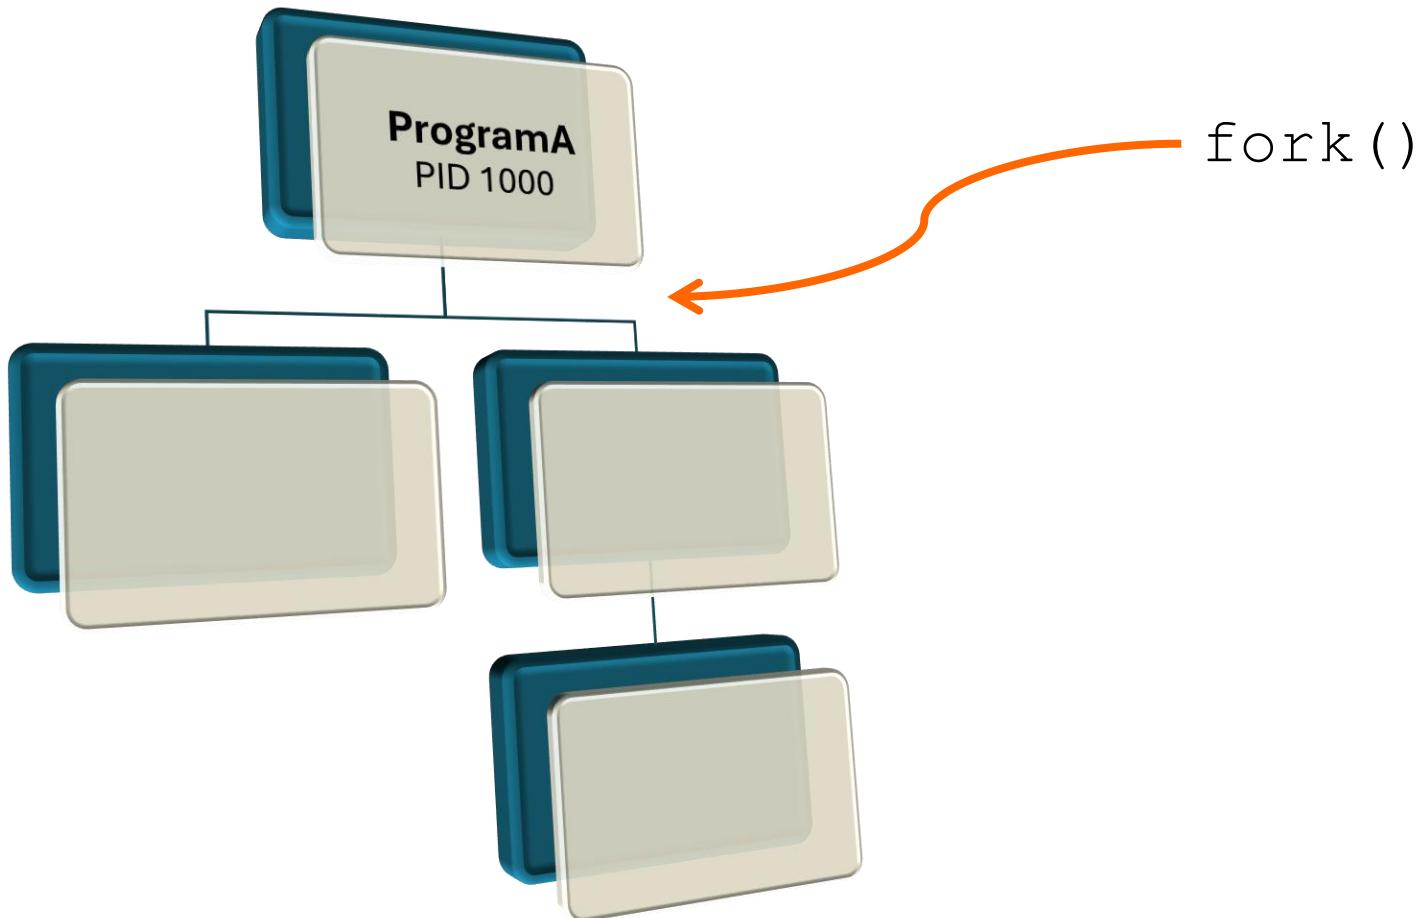


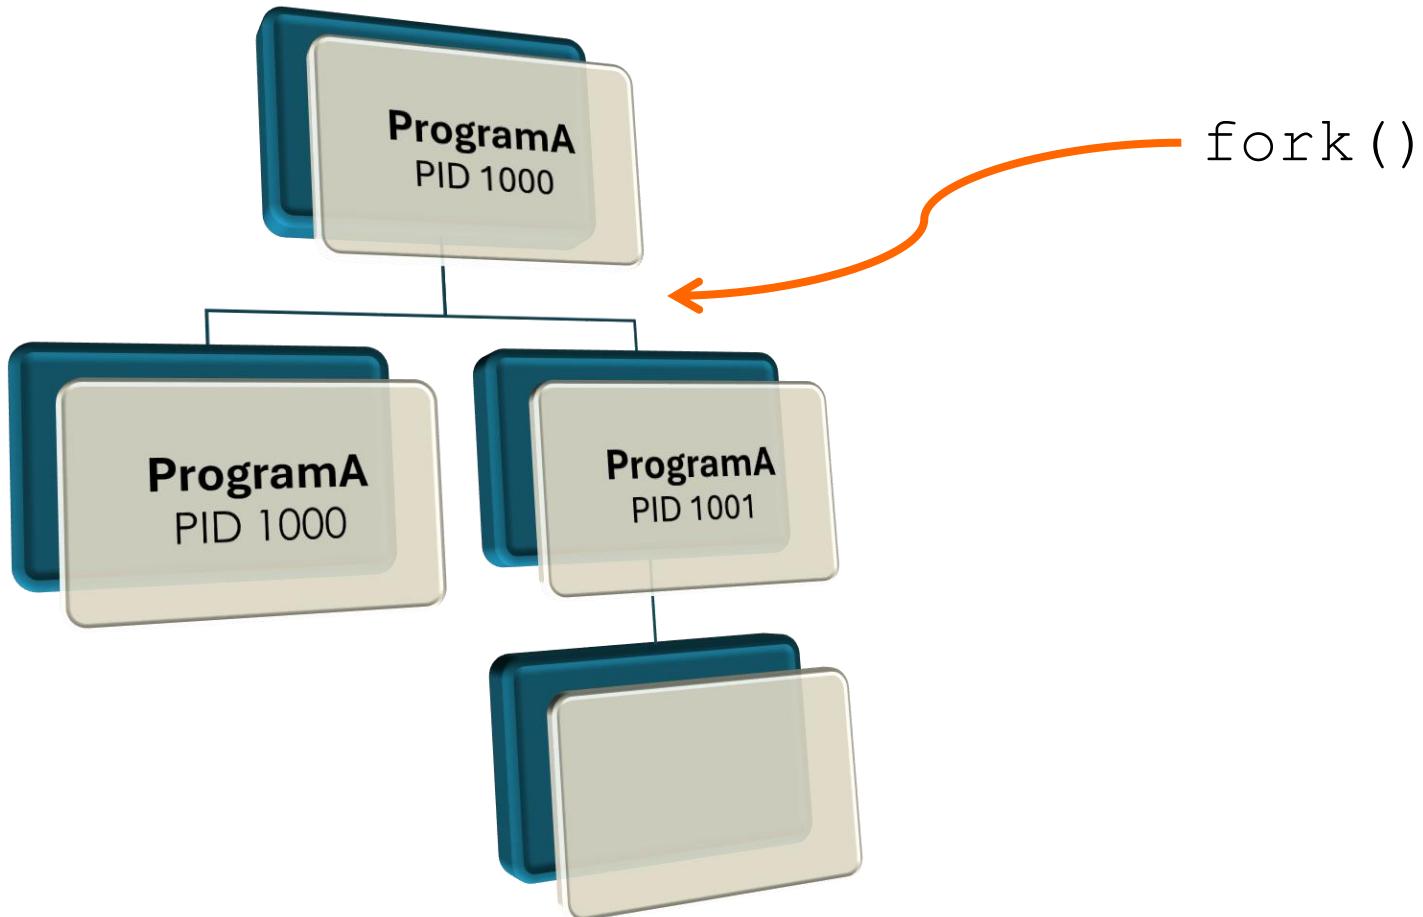


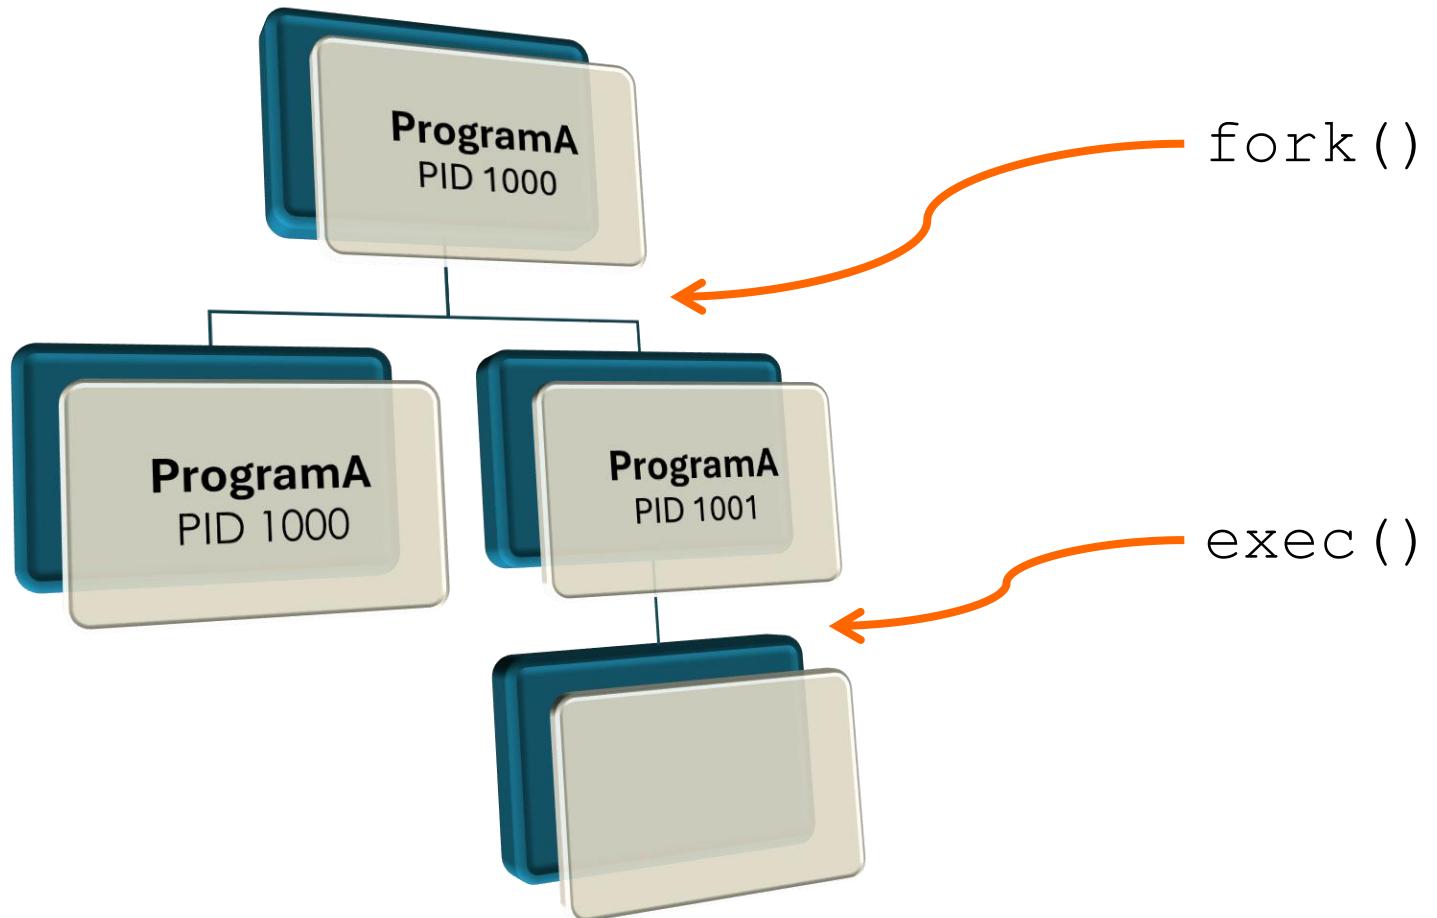
Review

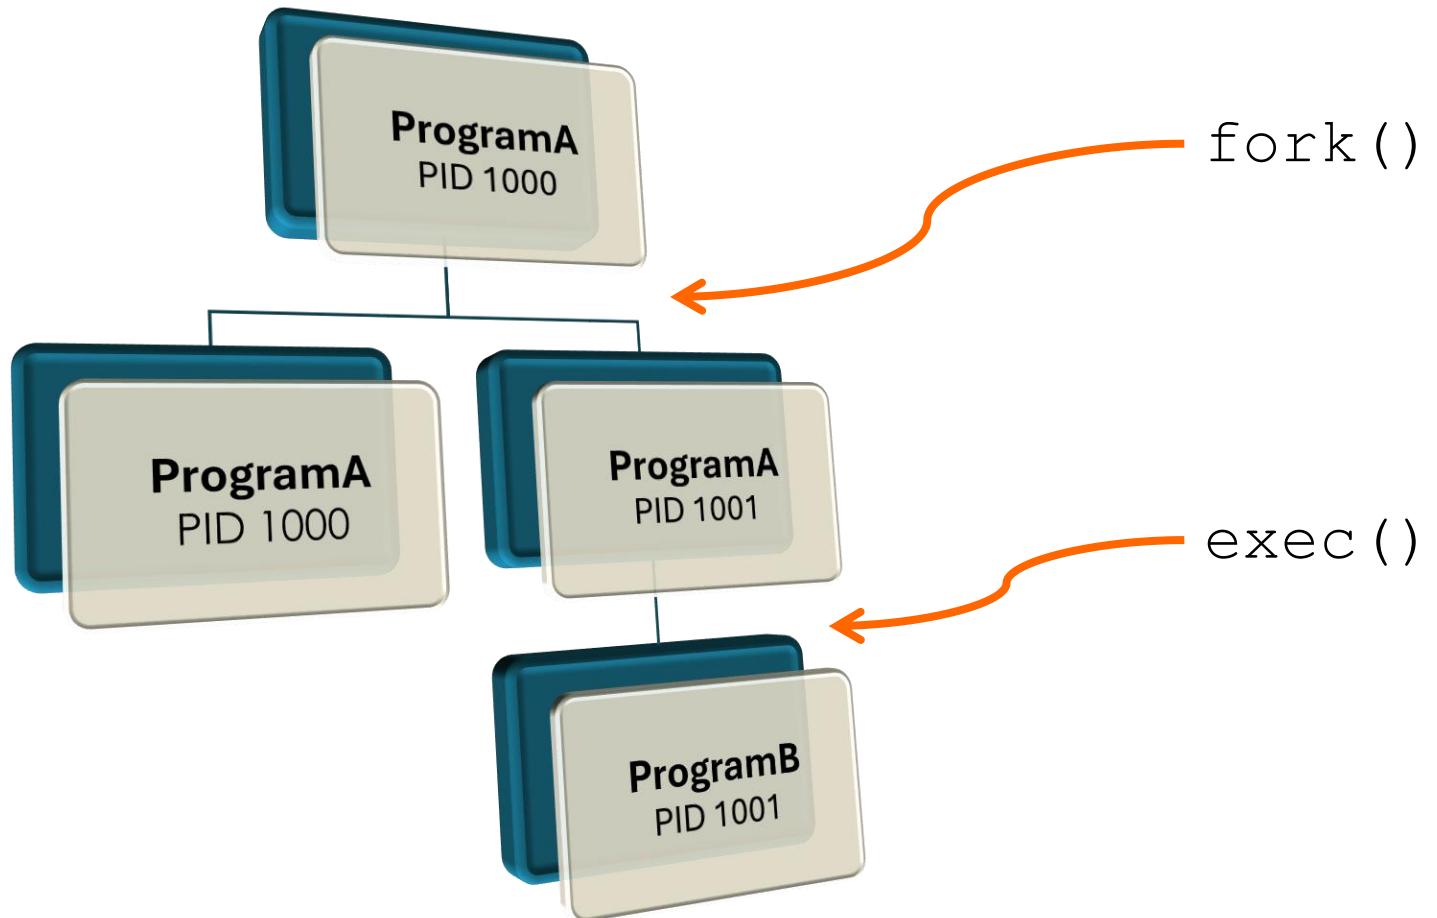
Processes, Threads, fork() and exec()













Questions?



UNIVERSITY OF
TORONTO